

A decorative graphic on the left side of the slide consisting of a network of thin, dark blue lines. These lines branch out and connect to small, empty circles, resembling a circuit board or a neural network diagram. The lines and circles are arranged in a way that suggests a flow or connection from the top left towards the bottom left.


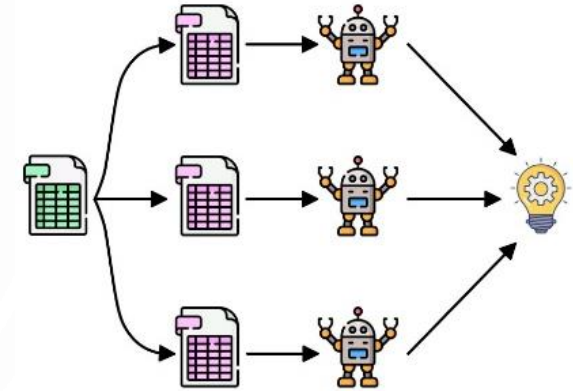
CSC 462: Machine Learning

7.5 Ensemble Learning

Dr. Sultan Alfarhood


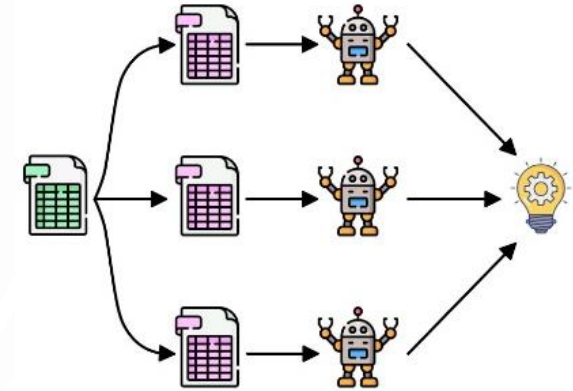
7.5 Ensemble Learning

- Ensemble learning **combines** several models to improve overall accuracy of ML algorithms
- Training a large number of low-accuracy models and then combining the predictions given by those weak models to obtain a high-accuracy meta-model



7.5 Ensemble Learning


- **Weak learners** are learning algorithms that cannot learn complex models
 - Typically **fast** at the training and at the prediction time
 - The most frequently used weak learner is a **decision tree** learning algorithm (stop splitting the training set after just a few iterations)
 - If the trees are not identical and each tree is at least slightly better than random guessing, then we can obtain high accuracy by combining a large number of such trees




7.5 Ensemble Learning

- Two most widely used and effective ensemble learning algorithms are
 - Random Forest
 - Gradient Boosting


Single Decision Tree



Gradient Boosted Trees

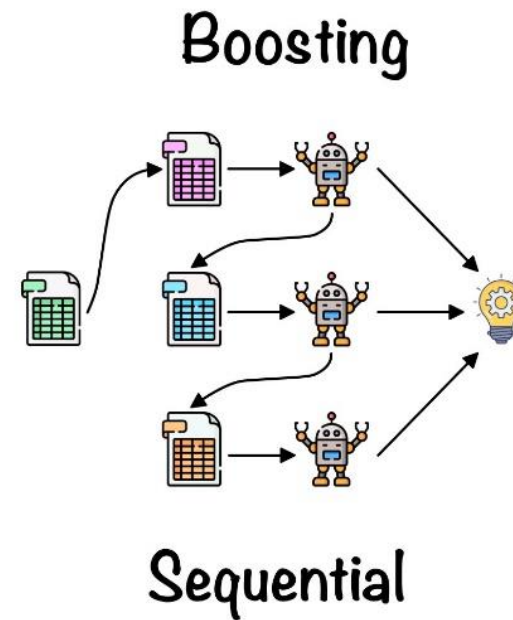
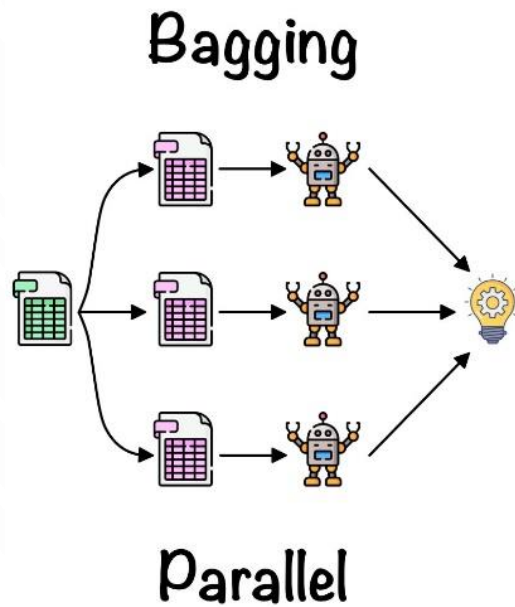


Random Forest




7.5 Ensemble Learning

- Ensemble learning paradigms
 - Bagging (Used in Random Forest)
 - Boosting (Used in Gradient Boosting)

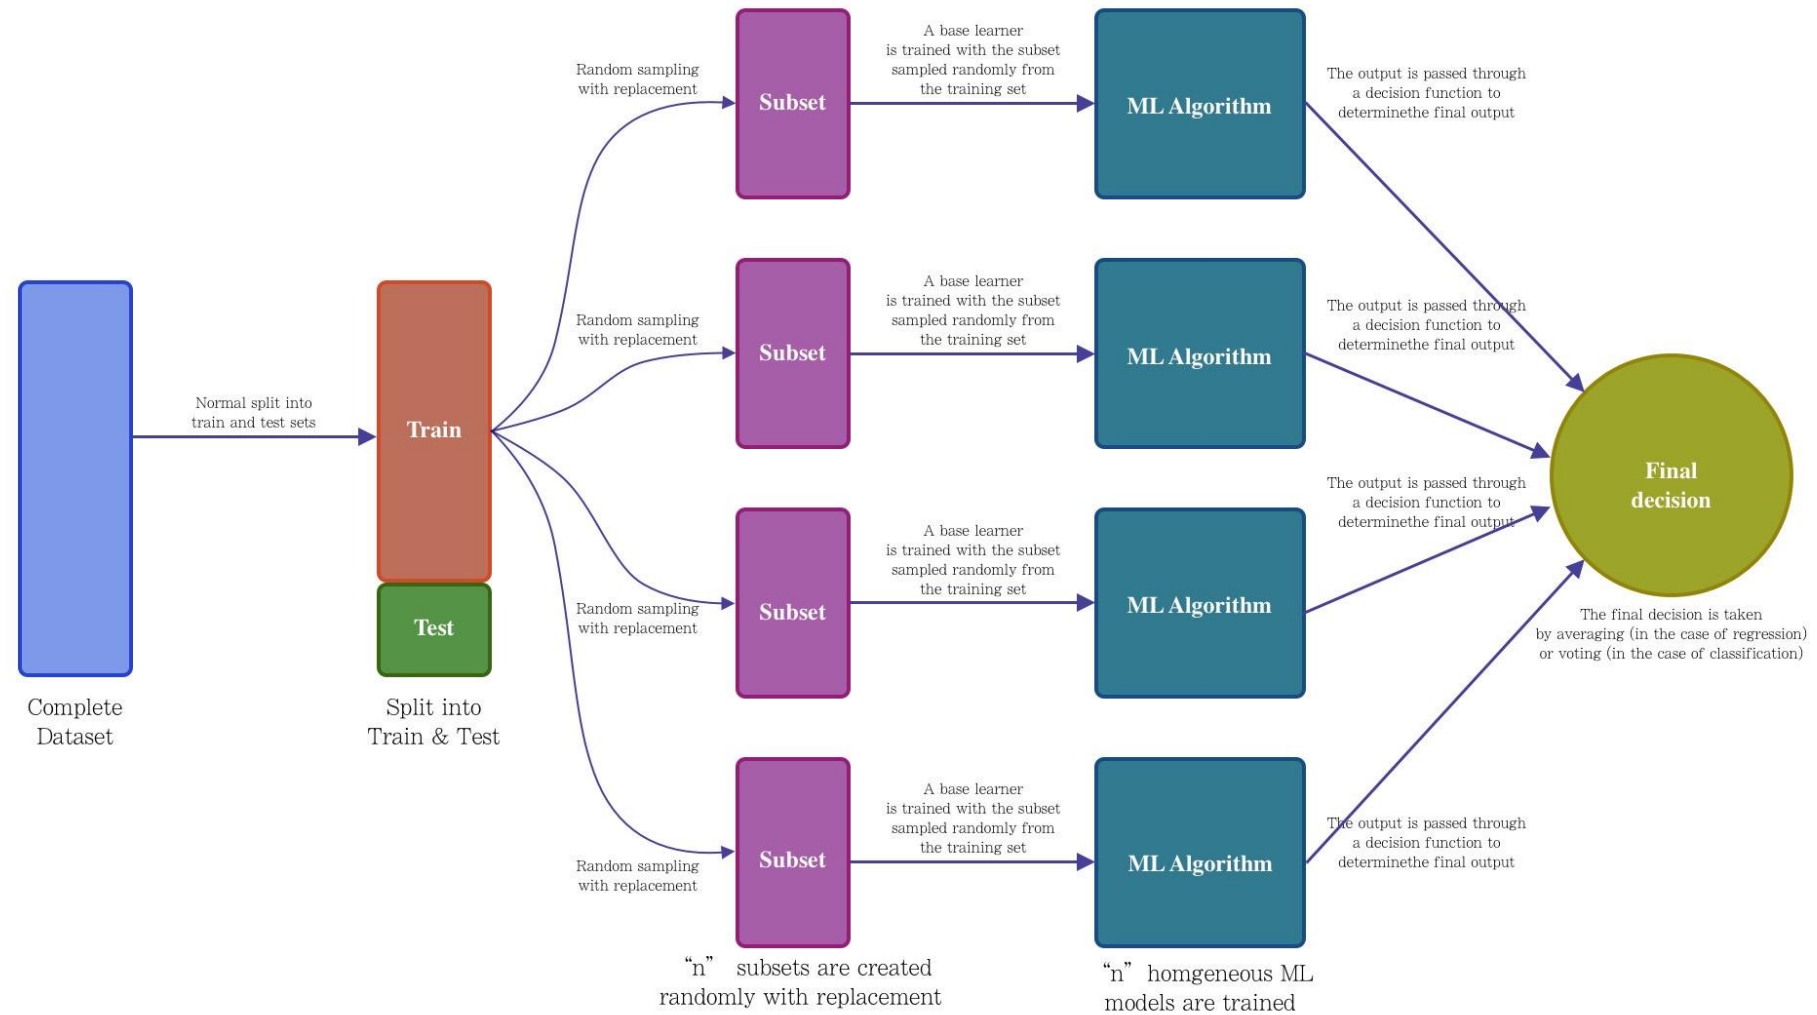


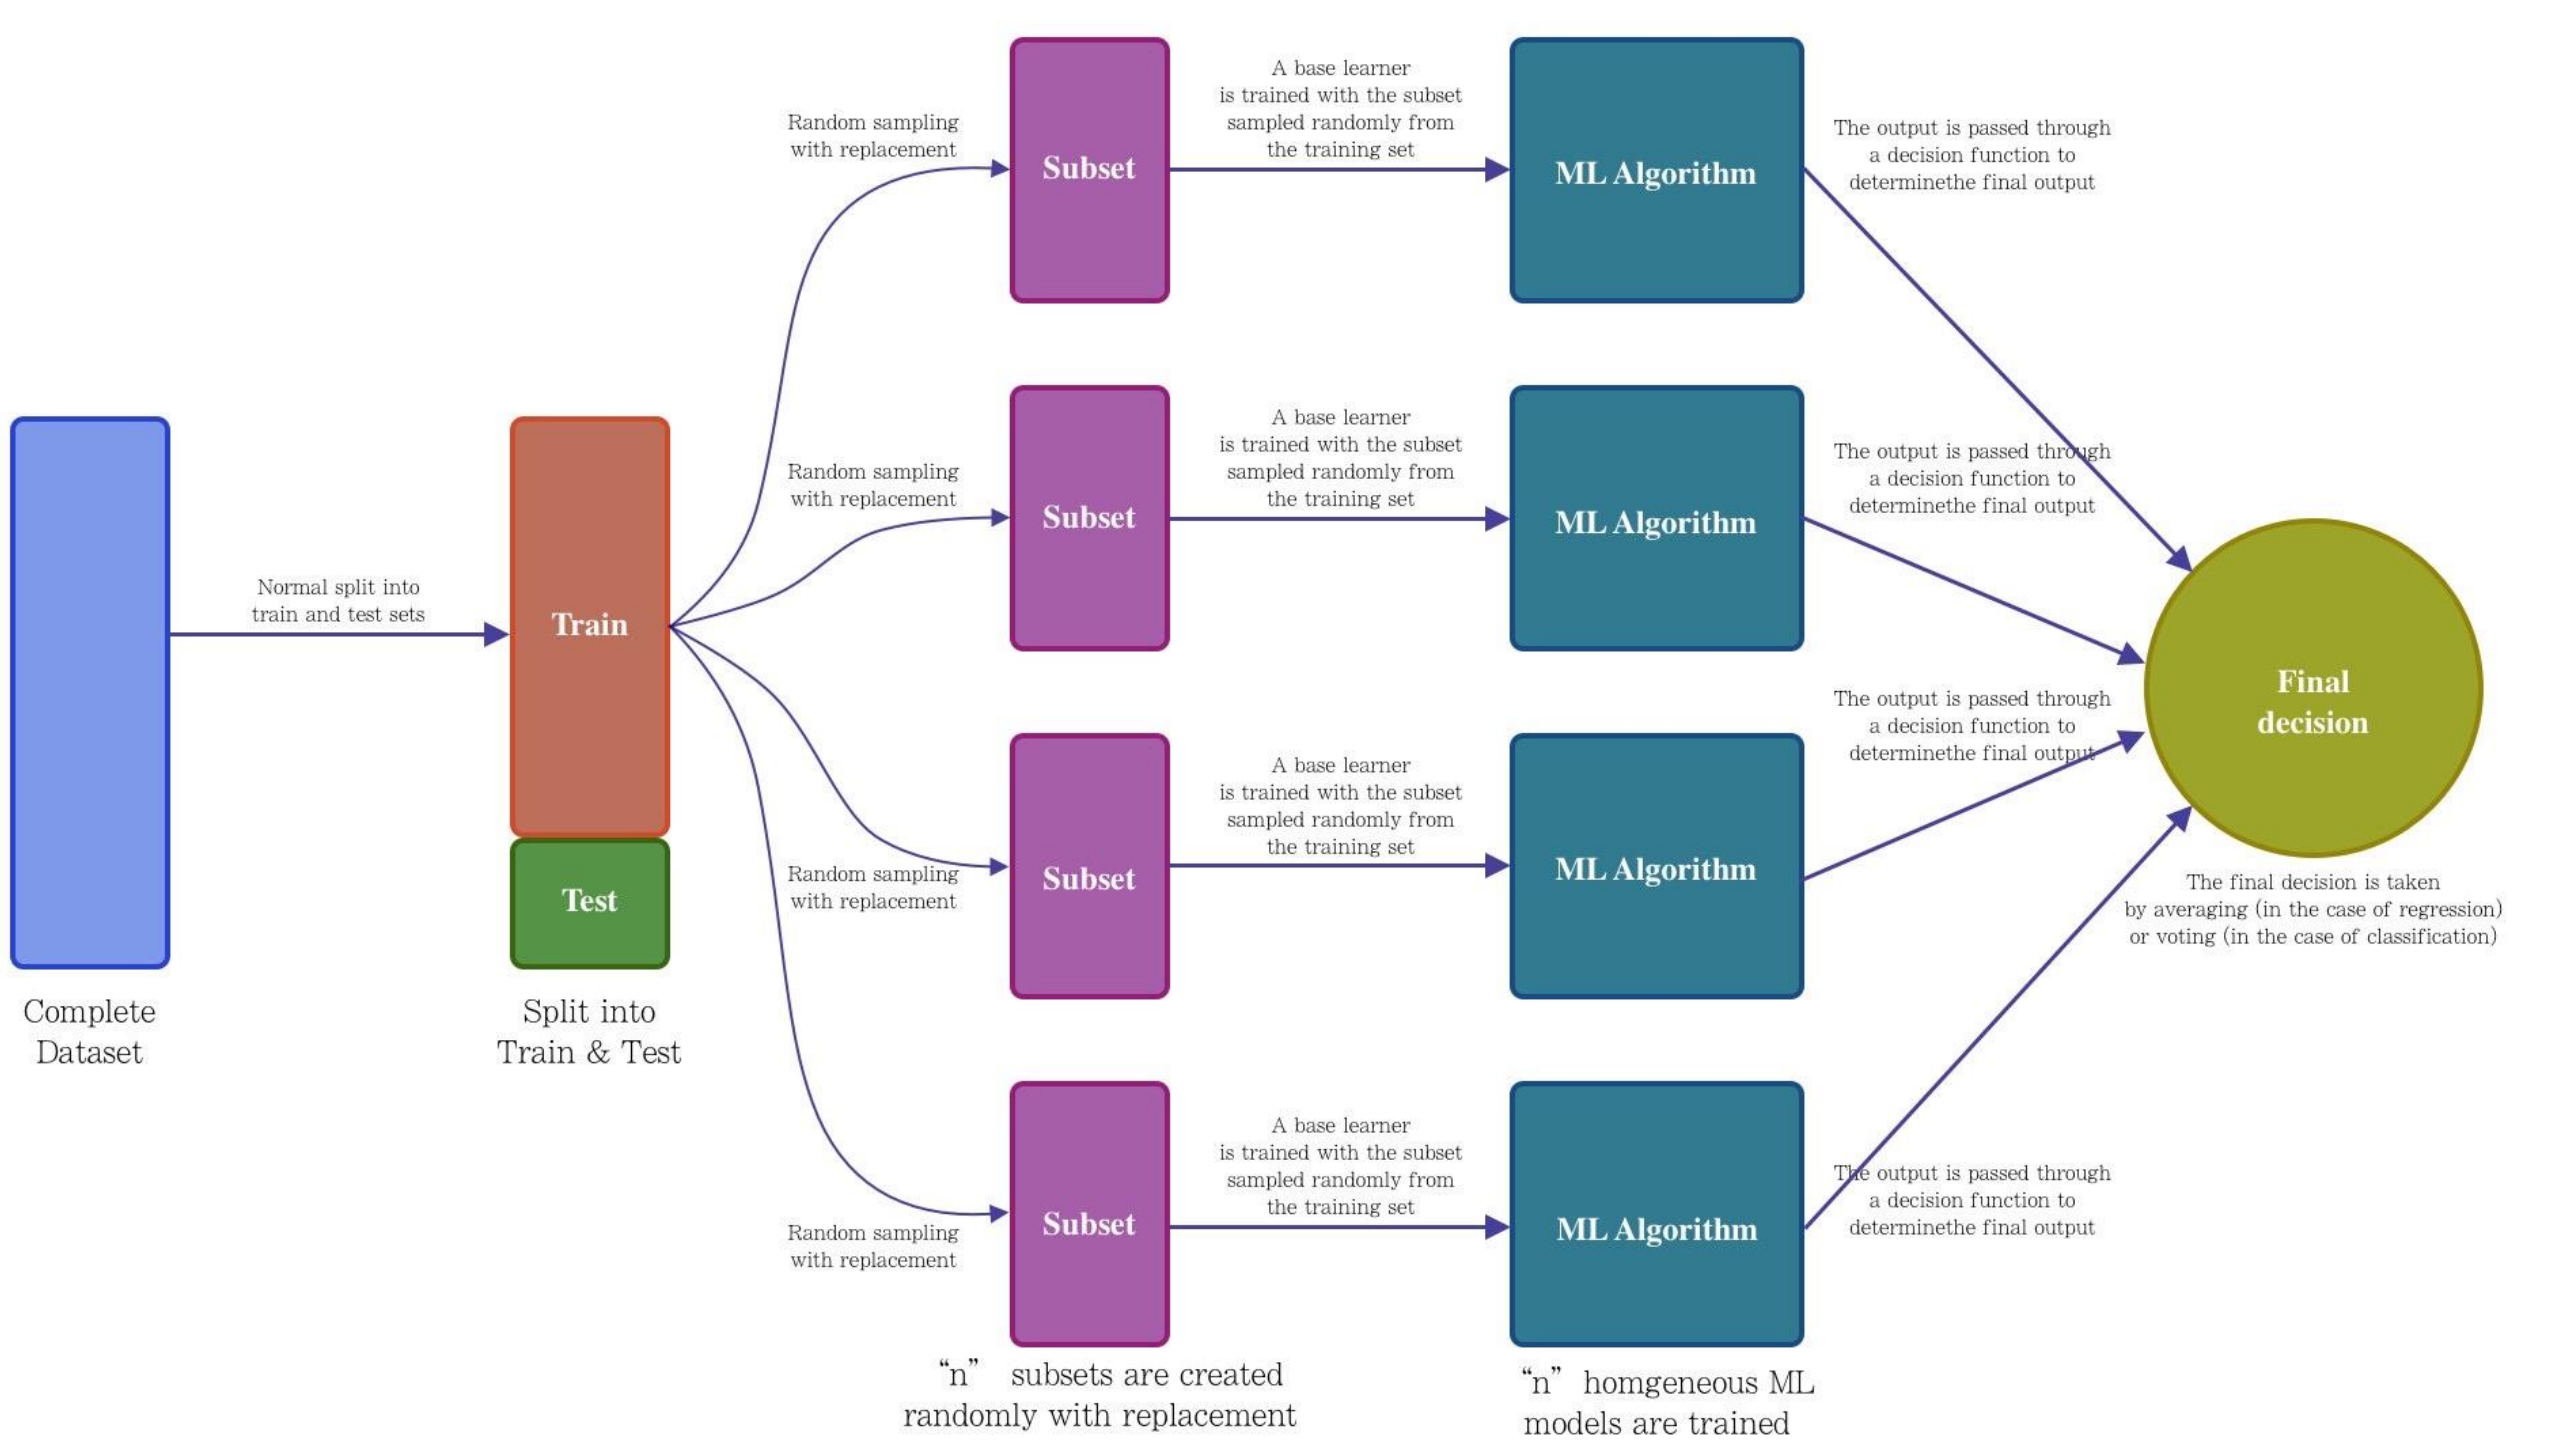
Bagging

1. Bagging consists of **creating many “copies”** of the training data
 - Each copy is slightly different from another
2. **Apply the weak learner** to each copy to obtain multiple weak models
3. **Combine** them



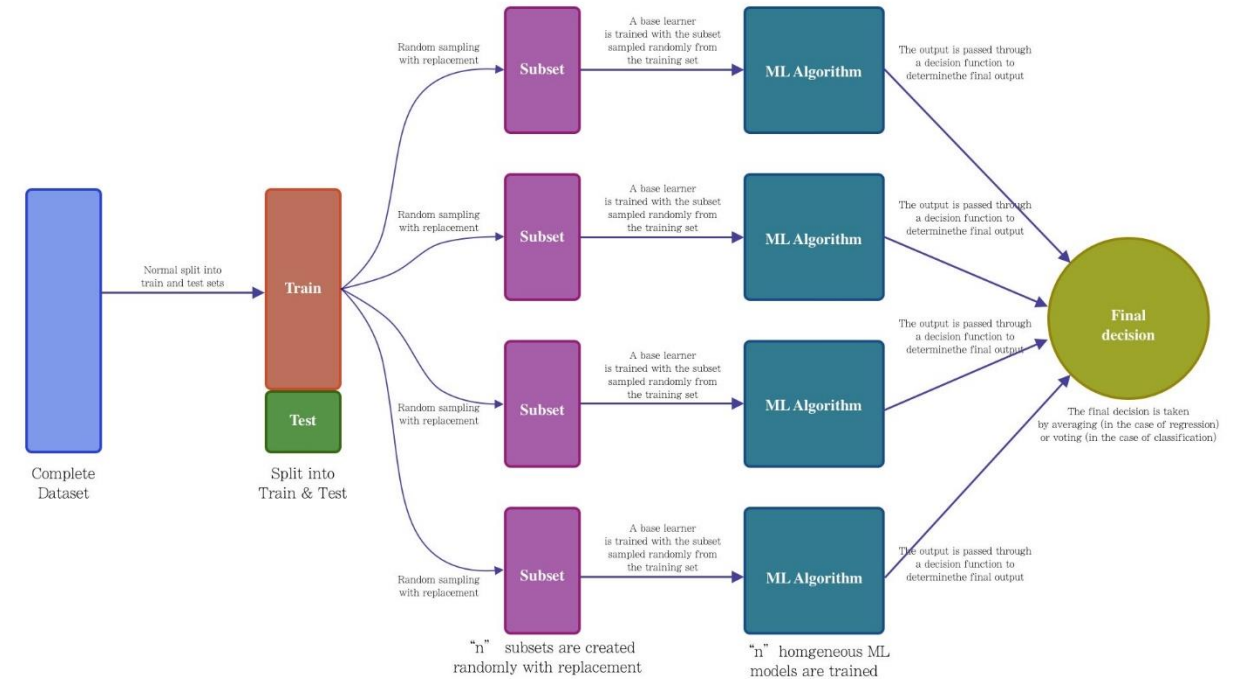
Random Forest





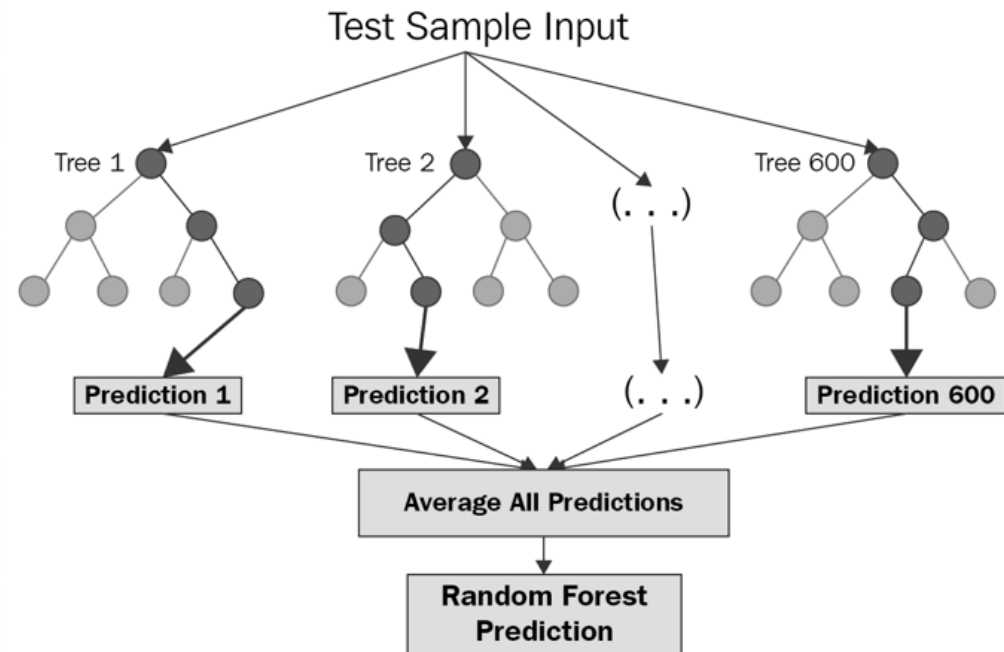
Random Forest

1. Given a training set, we create random samples of the training set and build a decision tree model using each sample as the training set
2. After training, we have N decision trees
3. The prediction for a new example x is obtained as the average of N predictions in the case of regression
 - Or by taking the majority vote in the case of classification



Random Forest

- The random forest algorithm is different from the vanilla bagging by using a modified tree learning algorithm that **inspects**, at each split in the learning process, a **random** subset of the features to **avoid the correlation** of the trees
- The most important **hyperparameters** to tune
 1. Number of trees
 2. Size of the random subset of the features to consider at each split



Random Forest (scikit-learn)

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification
X, y = make_classification(n_samples=1000, n_features=4,
                          n_informative=2, n_redundant=0,
                          random_state=0, shuffle=False)
clf = RandomForestClassifier(max_depth=2, random_state=0)
clf.fit(X, y)

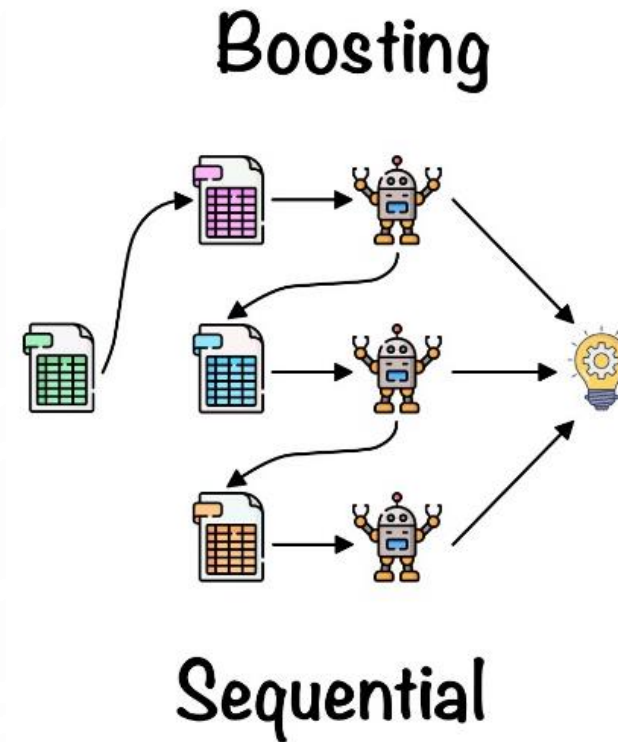
print(clf.predict([[0, 0, 0, 0]]))
```

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.datasets import make_regression
X, y = make_regression(n_features=4, n_informative=2,
                      random_state=0, shuffle=False)
regr = RandomForestRegressor(max_depth=2, random_state=0)
regr.fit(X, y)

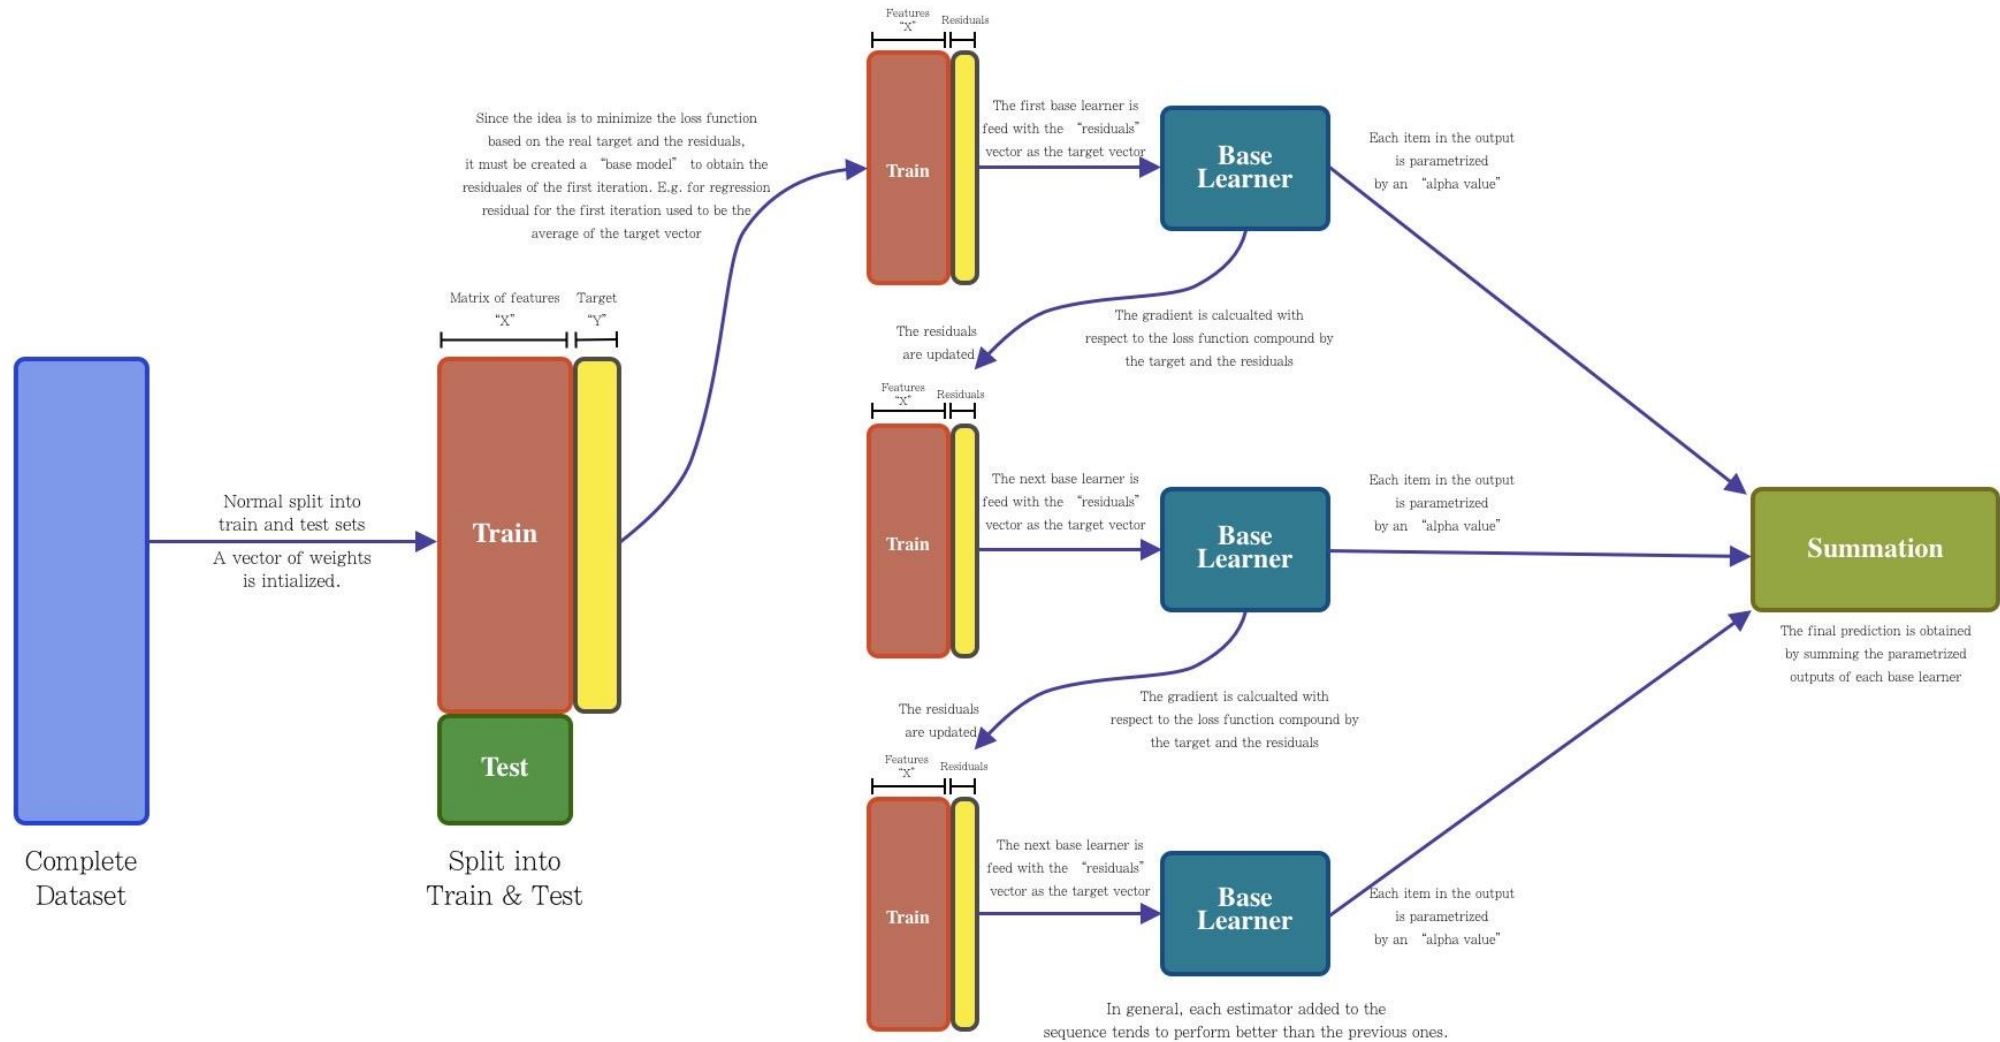
print(regr.predict([[0, 0, 0, 0]]))
```

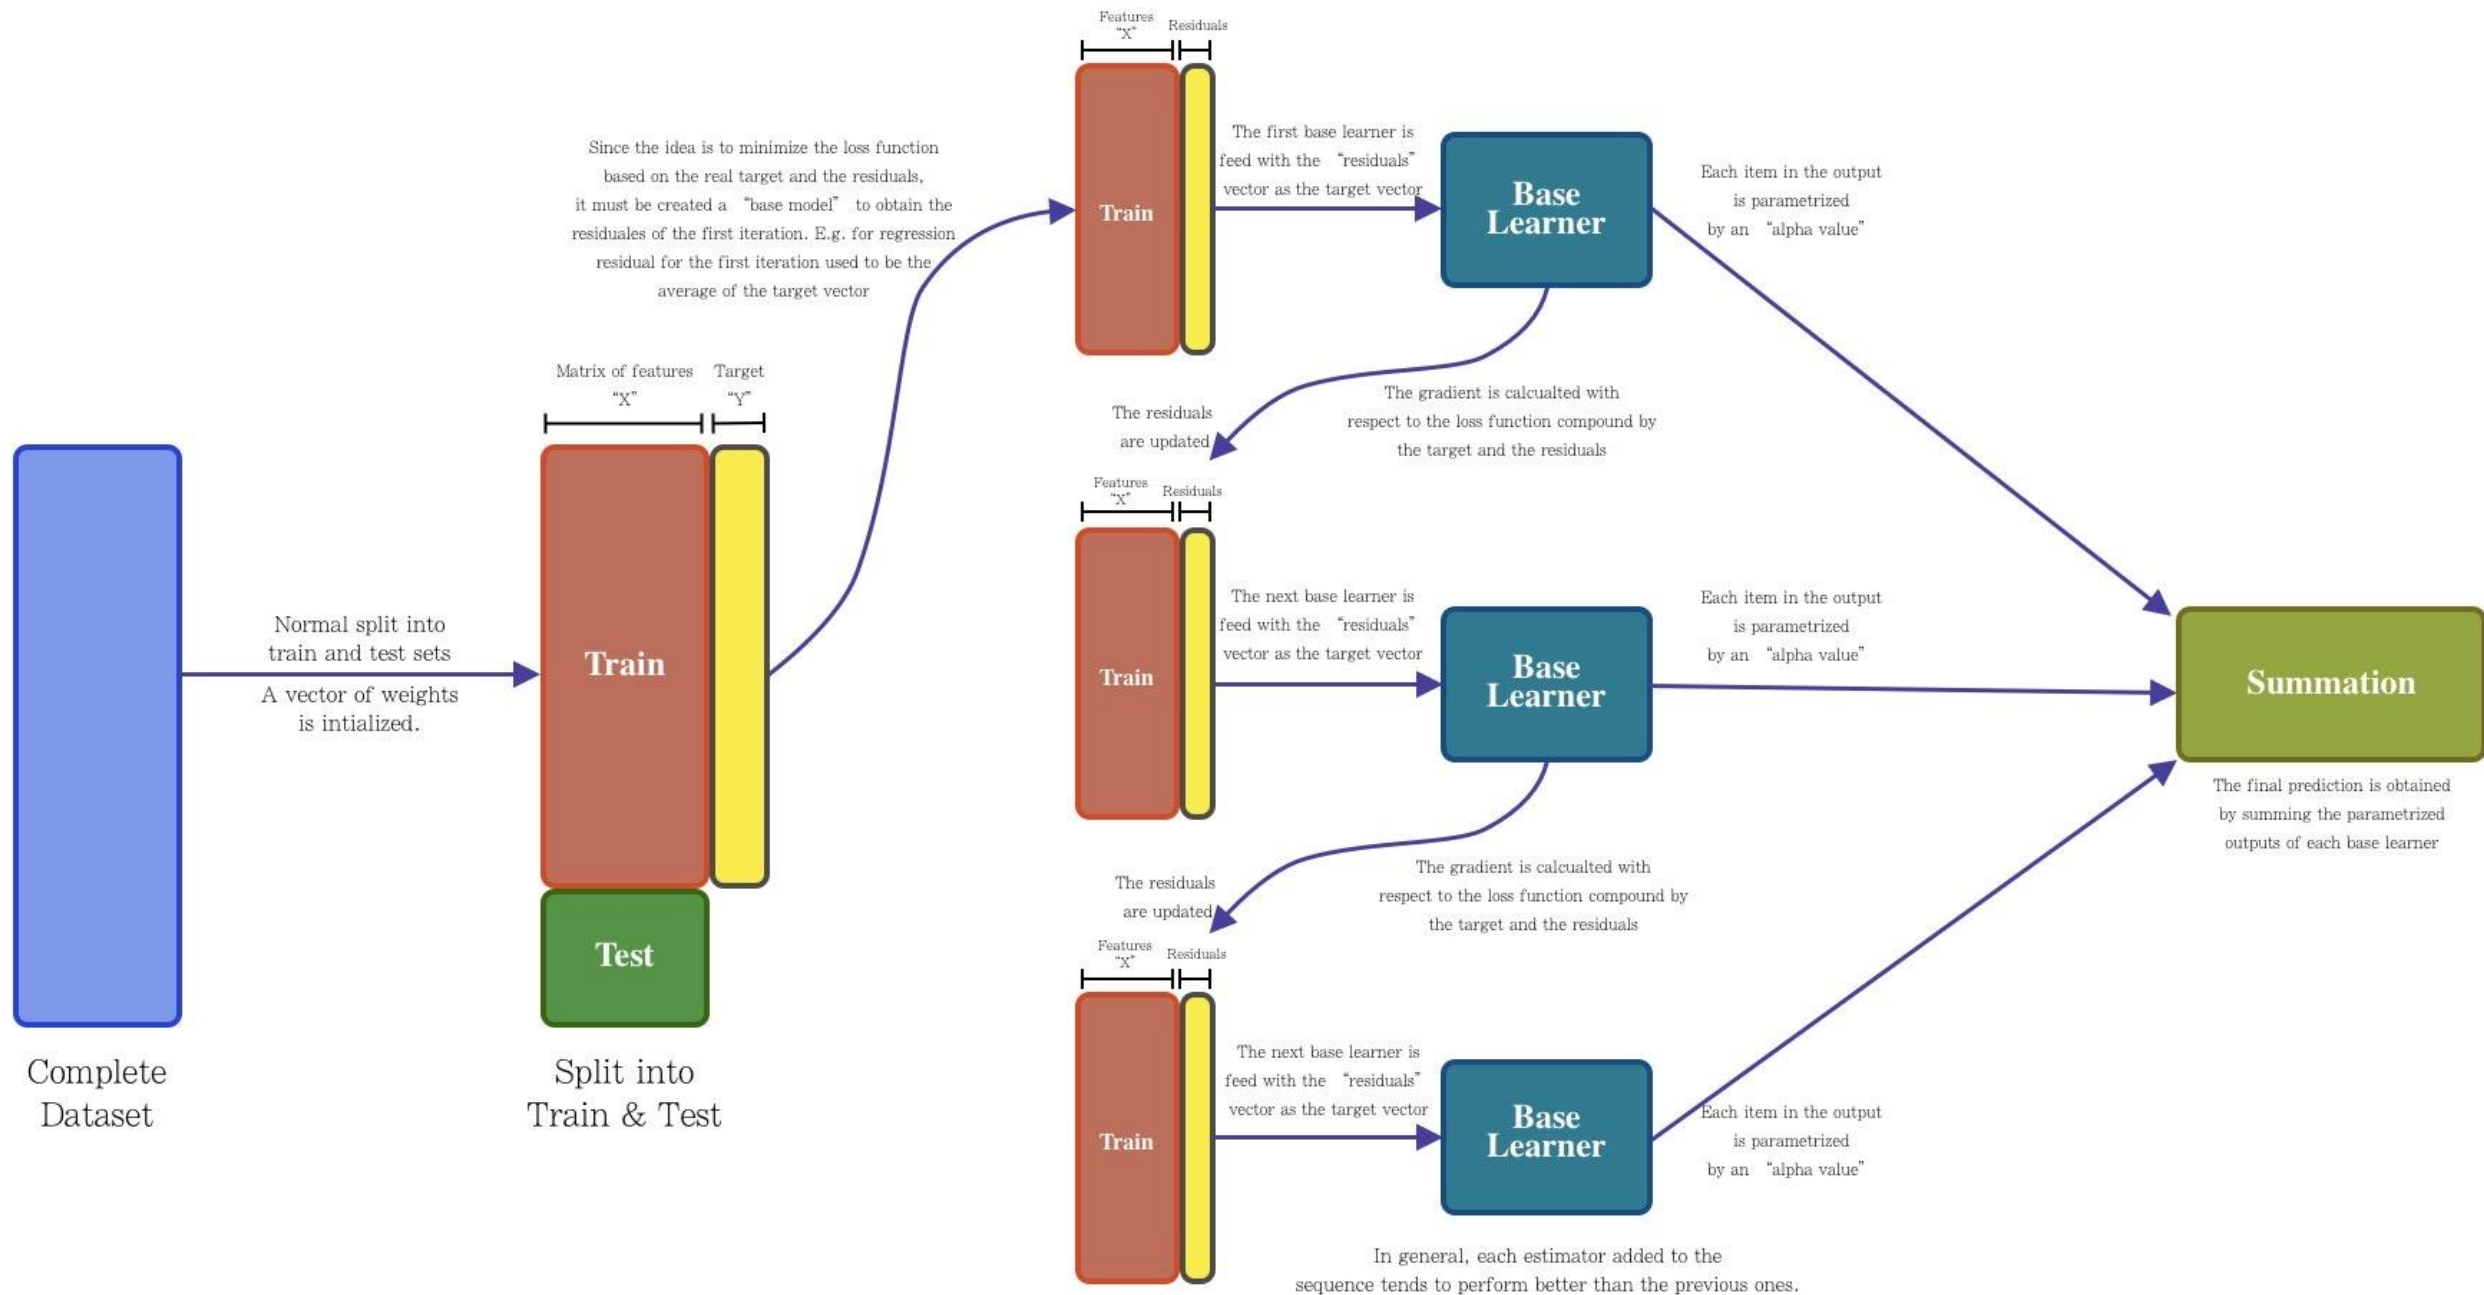
Boosting

- Boosting trains many models in a **gradual**, **additive** and **sequential** manner



Gradient Boosting





Gradient Boosting

1. Start with a constant model:

$$f = f_0 = \frac{\sum_{i=1}^N y_i}{N}$$

2. Modify labels of each example in our training set like follows:

$$\hat{y}_i \leftarrow y_i - f(x_i)$$

- where \hat{y}_i , called the **residual**, is the new label for example x_i

3. Use the modified training set, with residuals instead of original labels, to build a new decision tree model, f_1

- The boosting model is now defined as

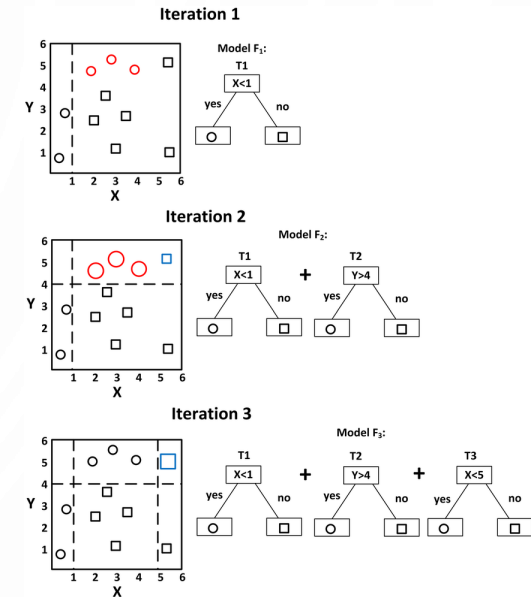
$$f = f_0 + \alpha f_1$$

- where α is the learning rate (a hyperparameter)

4. Recompute the residuals using equation in step 2 and replace the labels in the training data once again, train the new decision tree model f_2 , redefine the boosting model as

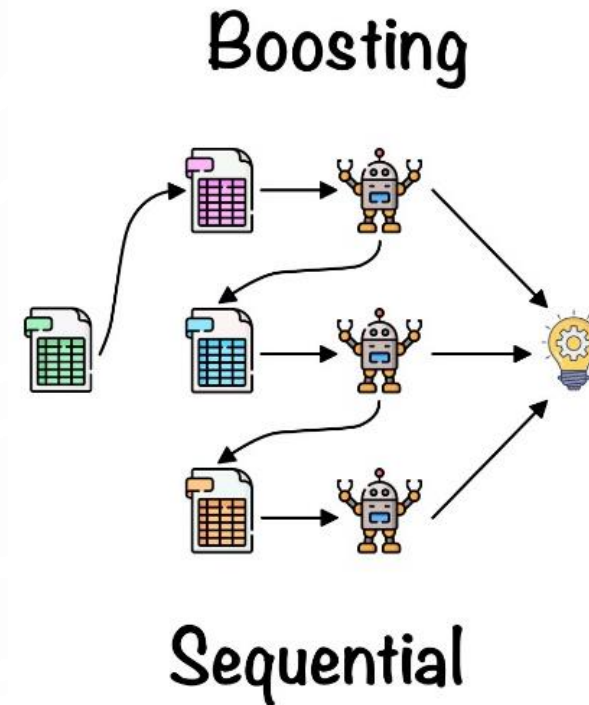
$$f = f_0 + \alpha f_1 + \alpha f_2$$

- And the process continues until the maximum of **M** (another hyperparameter) trees are combined




Gradient Boosting

- **Hyperparameters** to tune in gradient boosting:
 1. Number of trees
 2. Learning rate
 3. Depth of trees
- All the hyperparameters affect model accuracy
- The depth of trees also affects the speed of training and prediction: the shorter, the faster



Gradient Boosting

- Gradient boosting is one of the most powerful machine learning algorithms
 - It creates very **accurate** models
 - It is capable of **handling huge datasets** with millions of examples and features
- It usually outperforms random forest in accuracy but can be significantly **slower** in training because of its sequential nature



Gradient Boosting (scikit-learn)

```
from sklearn.datasets import make_hastie_10_2
from sklearn.ensemble import GradientBoostingClassifier
```

```
X, y = make_hastie_10_2(random_state=0)
X_train, X_test = X[:2000], X[2000:]
y_train, y_test = y[:2000], y[2000:]
```

```
clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,
                                max_depth=1, random_state=0).fit(X_train, y_train)
clf.score(X_test, y_test)
```

```
from sklearn.datasets import make_regression
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split
X, y = make_regression(random_state=0)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, random_state=0)
reg = GradientBoostingRegressor(random_state=0)
reg.fit(X_train, y_train)

reg.predict(X_test[1:2])

reg.score(X_test, y_test)
```

TensorFlow Decision Forests (TF-DF)

- TensorFlow Decision Forests (TF-DF) is a library for the training, evaluation, interpretation and inference of Decision Forest models.
- The two most popular DF training algorithms are Random Forests and Gradient Boosted Decision Trees.



TensorFlow Decision Forests (TF-DF)

https://colab.research.google.com/github/tensorflow/decision-forests/blob/main/documentation/tutorials/beginner_colab.ipynb

```
# Install TF-DF
!pip install tensorflow tensorflow_decision_forests

# Load TF-DF
import tensorflow_decision_forests as tfdf
import pandas as pd

# Load a dataset in a Pandas dataframe.
train_df = pd.read_csv("project/train.csv")
test_df = pd.read_csv("project/test.csv")

# Convert the dataset into a TensorFlow dataset.
train_ds = tfdf.keras.pd_dataframe_to_tf_dataset(train_df, label="my_label")
test_ds = tfdf.keras.pd_dataframe_to_tf_dataset(test_df, label="my_label")

# Train a Random Forest model.
model = tfdf.keras.RandomForestModel()
model.fit(train_ds)

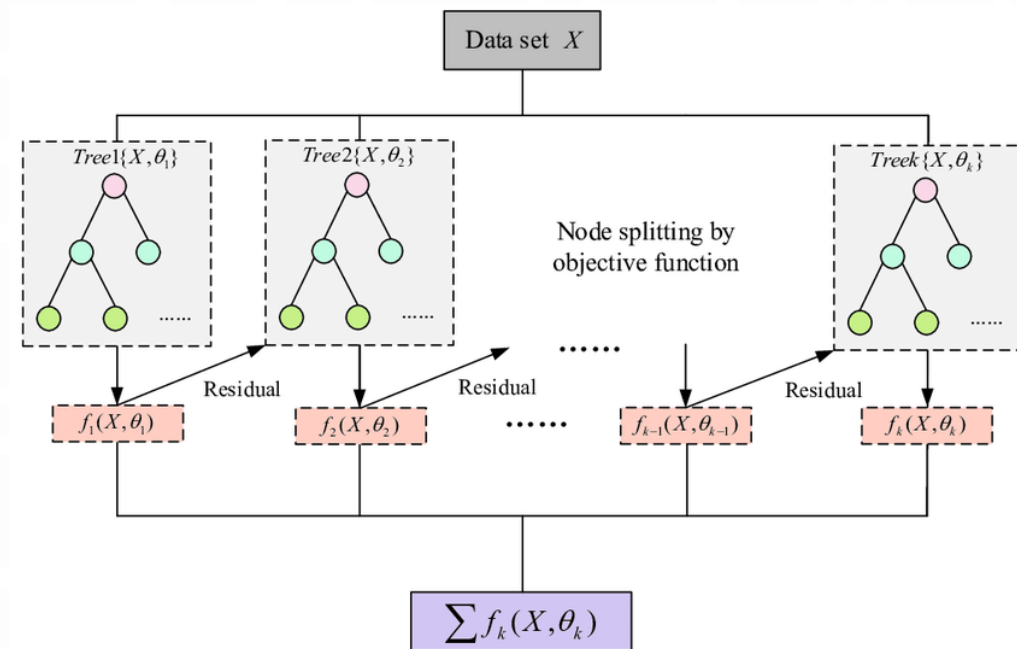
# Summary of the model structure.
model.summary()

# Evaluate the model.
model.evaluate(test_ds)

# Export the model to a SavedModel.
model.save("project/model")
```

XGBoost

- XGBoost, which stands for Extreme Gradient Boosting, is a scalable, distributed gradient-boosted decision tree machine learning library.
- It provides **parallel tree boosting** and is the **leading** machine learning library for regression, classification, and ranking problems.



XGBoost

```
from numpy import loadtxt
from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# load data
dataset = loadtxt('pima-indians-diabetes.csv', delimiter=",")

# split data into X and y
X = dataset[:,0:8]
Y = dataset[:,8]

# split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, random_state=42)

# fit model no training data
model = XGBClassifier()
model.fit(X_train, y_train)

# make predictions for test data
y_pred = model.predict(X_test)

# Predictions made by XGBoost are probabilities. Convert them to binary class values by rounding them to 0 or 1.
predictions = [round(value) for value in y_pred]

# evaluate predictions
accuracy = accuracy_score(y_test, predictions)
print("Accuracy: %.2f%%" % (accuracy * 100.0))
```