

SOLVING PROBLEMS BY SEARCHING

Outline

- ❖ Problem-Solving Agents
 - Search problems and solutions
 - Formulating problems
- ❖ Example Problems
 - Standardized problems
 - Real-world problems
- ❖ Search Algorithms
 - Best-first search
 - Search data structures
 - Redundant paths
 - Measuring problem-solving performance
- ❖ Uninformed Search Strategies
 - Breadth-first search
 - Dijkstra's algorithm or uniform-cost search
 - Depth-first search and the problem of memory

Problem-Solving Agents

- ❑ Problem-solving agent: An agent that
 - plans ahead: to consider a sequence of actions that form a path to a goal state
 - when correct action to take is not immediately obvious
 - computational process it undertakes is called search
- ❑ For sake of simplicity, environment is:
 - episodic, single agent, fully observable, deterministic, static, discrete, and known
- ❑ Two types of algorithms:
 - informed algorithms: in which agent can estimate how far it is from the goal
 - uninformed algorithms: where no such estimate is available.

Problem-Solving Agents

- ❑ Agent is currently in city A, trying to reach destination Z
 - Observes that there are three roads leading out of A: one toward B, one to C, and one to D
 - None of these are the goal
 - Unless agent is familiar w/ geography of its world, it will not know which road to follow
 - If no additional info. (environment is unknown) agent can do no better than random actions
- ❑ Assume agents always have access to info. about world (map)
 - follow four-phase problem-solving process:
 - **Goal formulation:** agent adopts goal of reaching Z.
 - Goals organize behavior by limiting objectives → actions to be considered
 - **Problem formulation:** agent devises a description of states and actions necessary to reach goal
 - abstract model of relevant part of world
 - one good model: consider actions of traveling from one city to an adjacent city → current city
 - **Search:** before taking any action in real world, agent simulates sequences of actions in its model
 - searching until it finds solution: sequence of actions that reaches goal.
 - might have to simulate multiple sequences that do not reach goal, or find that no solution is possible
 - **Execution:** agent can now execute actions in solution, one at a time

Problem-Solving Agents

- ❑ In fully observable/deterministic/known environment: solution is fixed actions sequence
 - If model is correct
 - once agent has found a solution, it can ignore its percepts while executing actions
 - Open-loop: because solution is guaranteed to lead to goal
 - If there is chance model is incorrect, or environment is nondeterministic
 - agent would be safer using:
 - Closed-loop approach that monitors the percepts.
- ❑ In partially observable/non-deterministic environments:
 - solution would be branching strategy that recommends different future actions depending on percepts
 - Ex. agent might plan to drive from A to B but might need contingency plan in case:
 - it arrives in Z by accident
 - or finds a sign saying “Road Closed”

Problem-Solving Agents

Search Problems and Solutions

□ Search Problem:

- set of possible **states** of environment: **state space**, can be represented as graph vertices
- **initial state** that agent starts in
- set of **goal states**
 - can be empty, multiple, property that applies to many/infinite states
 - Ex. in vacuum-cleaner world, goal might be to have no dirt in any location
 - specify IS-GOAL method for problem
- **actions** available to agent, can be the graph directed edges
 - Given state s, ACTIONS(s) returns a finite set of actions that can be executed in s.
 - We say that each of these actions is **applicable** in s
 - **path**: sequence of actions
 - **solution**: path from initial state to a goal state.
- **transition model**: describes what each action does
 - RESULT(s, a) returns state that results from doing action a in state s
- **action cost function**: ACTION-COST(s, a, s₀) in programming or c(s, a, s₀) in math
 - gives numeric cost of applying action a in state s to reach state s₀
 - should use cost function that reflects its own performance measure
 - Ex. in route-finding agents, cost of action might be distance in miles, or time it takes to complete action.
 - **optimal solution** has lowest path cost among all solutions

Problem-Solving Agents

Search problems and solutions

❑ Notes:

- Real world is absurdly complex → state space must be abstracted for problem solving
 - (Abstract) state = set of real states
 - (Abstract) action = complex combination of real actions
 - (Abstract) solution = set of real paths that are solutions in the real world
 - Each abstract action should be "easier" than the original problem!

Example Problems

❑ Standardized problem:

- intended to illustrate/exercise various problem-solving methods
- It can be given a concise, exact description
→ suitable as benchmark for Real-world problem researchers to compare performance of algorithms.

❑ Real-world problem:

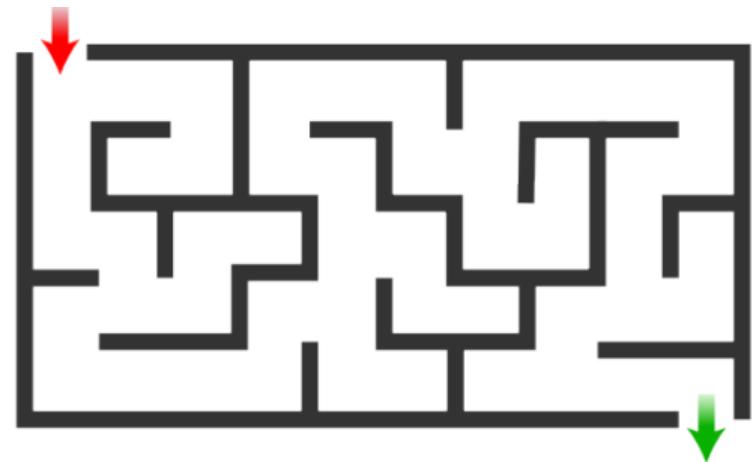
- one whose solutions people actually use,
- whose formulation is idiosyncratic, not standardized
because, for example, each robot has different sensors that produce different data.

Example Problems

Standardized problems

□ Grid world problem

- 2D/3D rectangular array of square cells in which agents can move from cell to cell
- agent can move to any obstacle-free adjacent cell horizontally or vertically, in some problems diagonally.
- cells can contain:
 - objects, which agent can pick up/push, or otherwise act upon
 - wall or other impassible obstacle in a cell prevents an agent from moving into that cell.

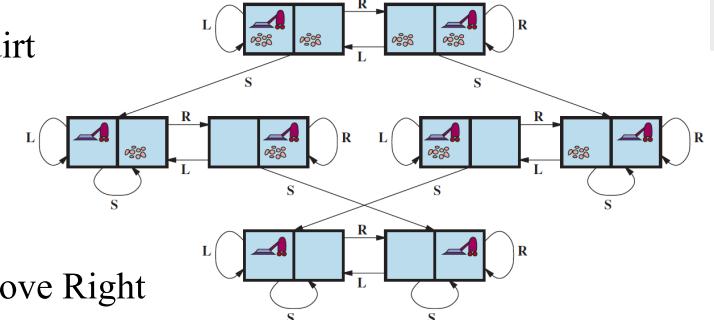


Example Problems

Standardized problems: Grid world problem

❑ Vacuum world as a grid world problem:

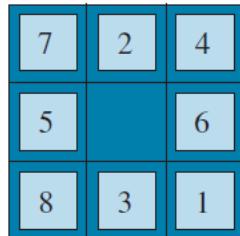
- **States:** which objects are in which cells: objects are agent and dirt
 - simple two-cell version, agent can be in either
 - each cell can either contain dirt or not
 - → there are $2 \times 2 \times 2 = 8$
 - In general, a vacuum environment with n cells has $n \cdot 2^n$ states.
- **Initial state:** Any state can be designated as the initial state.
- **Actions:** In two-cell world we defined: Suck, move Left, and move Right
 - In 2D multi-cell world we could add Upward and Downward
 - or switch to egocentric relative actions: Forward, Backward, TurnRight, and TurnLeft
- **Transition model:**
 - Suck removes any dirt from agent's cell: $([X,Y], \text{Suck}) \rightarrow [X, \text{CLEAN}]$
 - Forward moves agent ahead 1 cell in direction it is facing, unless it hits a wall, in which case action has no effect
 - Backward moves agent in opposite direction
 - TurnRight and TurnLeft change direction it is facing by 90
- **Goal states:** states in which every cell is clean.
- **Action cost:** Each action costs 1.



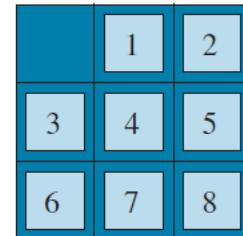
Example Problems

Standardized problems: Grid world problem

□ Sliding-tile puzzle as a grid world problem:

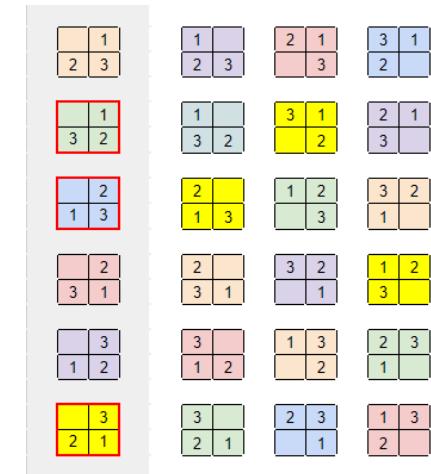


Start State



Goal State

- **States:** state description specifies location of each tile
- **Initial state:** Any state can be designated as initial state
 - parity property partitions state space
→ any given goal can be reached from exactly half of possible initial states
- **Actions:** tile slides (physical world), simpler to move blank Left,Right,Up,Down
 - If blank is at an edge or corner then not all actions will be applicable.
- **Transition model:** Maps state/action to resulting state
- **Goal state:** although any state could be goal, typically state w/ numbers in order
- **Action cost:** Each action costs 1.



- Configurations w/ same color obtained by rotation
- States w/ red border are unsolvable

Example Problems

Standardized problems

□ Knuth conjecture problem

- Starting with number 4, sequence of square root/floor/factorial operations can reach any positive integer
 - Ex. we can reach 5 from 4 as follows:
- state space is infinite: for any integer >2 factorial operator will always yield a larger integer
- problem is interesting because it explores very large numbers
 - shortest path to 5 goes through $(4!)! = 620,448,401,733,239,439,360,000$
- **States:** Positive real numbers.
- **Initial state:** 4.
- **Actions:** Apply square root, floor, or factorial operation (factorial for integers only).
- **Transition model:** As given by mathematical definitions of operations.
- **Goal state:** desired positive integer.
- **Action cost:** Each action costs 1.

$$\lfloor \sqrt{\sqrt{\sqrt{\sqrt{(4!)!}}}} \rfloor = 5$$

Example Problems

Real-world problems

❑ Route-finding problem: Ex. airline travel problems

- **States:** Each state includes a location , current time.
 - cost of action (flight segment) depends on: previous segments, their fare bases, domestic/international
 - → state must record extra information about these “historical” aspects.
- **Initial state:** user’s home airport.
- **Actions:** Take any flight from current location, in any seat class, leaving after current time, leaving enough time for within-airport transfer if needed.
- **Transition model:** state resulting from taking flight will have flight’s destination as new location and flight’s arrival time as new time.
- **Goal state:** destination city. Sometimes goal can be more complex, such as “arrive at destination on a nonstop flight.”
- **Action cost:** combination of monetary cost, waiting time, flight time, customs/immigration procedures, seat quality, time of day, type of airplane, frequent-flyer reward points, and so on.

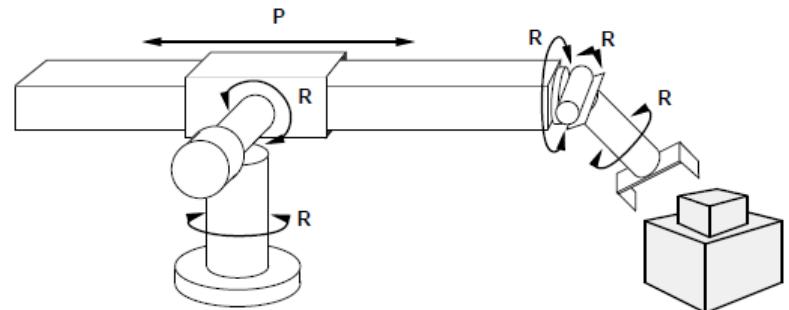
Example Problems

Real-world problems

Robot navigation problem: generalization of route-finding problem

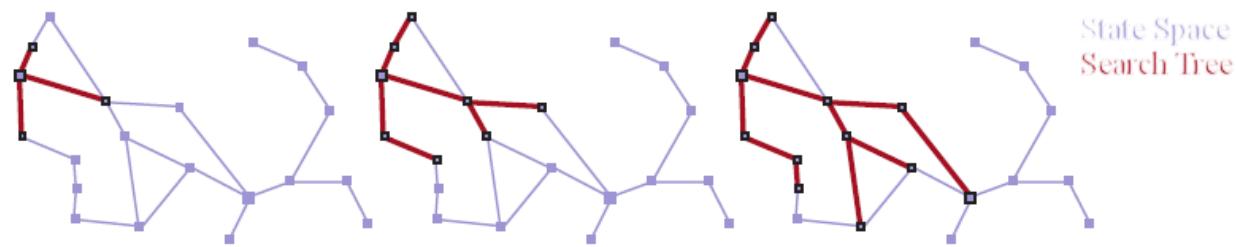
- roam around making its own paths vs following distinct paths
 - circular robot moving on flat surface → space is essentially 2D
 - robot has controlled arms/legs
 - search space becomes many-dimensional—one dimension for each joint angle
- Advanced techniques are required just to make essentially continuous search space finite
- complexity + errors in sensor readings/motor controls, partial observability, other agents alter environment

- **States:** real-valued coordinates of robot joint angles , parts of the object to be assembled
- **Actions:** continuous motions of robot joints.
- **Goal state:** complete assembly with no robot included!
- **Action cost:** time to execute



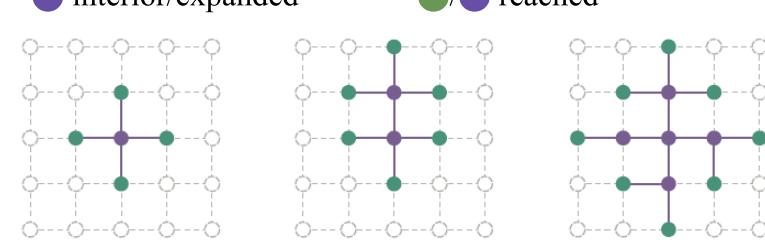
Search Algorithms

- ❑ Search algorithm takes search problem as input , returns solution / failure indication
- ❑ Algorithms that superimpose search tree over state space graph
 - forming various paths from initial state, trying to find path that reaches a goal state
 - each node in search tree corresponds to state in state space
 - edges in search tree correspond to actions
 - root of tree corresponds to initial state of problem
- ❑ Distinction between state space and search tree
 - state space describes set of states in world, and actions that allow transitions from state to another
 - search tree describes paths between these states, reaching towards goal
 - search tree may have multiple paths to (and thus multiple nodes for) any given state, but each node in tree has a unique path back to root



Search Algorithms

- ❑ Starting from root, expand node, by considering available ACTIONS for that state
 - use RESULT function to see where those actions lead to, generate child/successor node for each state
 - choose which of child nodes to consider next
 - search: following up one option now and putting others aside for later
- ❑ Reached states: states that have had nodes generated for them
- ❑ Frontier: children of child node that have not been expanded yet
 - frontier separates two regions of state-space graph:
 - interior region where every state has been expanded
 - exterior region of states that have not yet been reached
- ❑ Ex:
 - exterior/unreached
 - frontier/fringe
 - interior/expanded
 - / ● reached
 1. just root has been expanded
 2. top frontier node is expanded
 3. remaining successors of root are expanded in clockwise order



Search Algorithms

Best-first search

- Choose node n w/ min. value of some evaluation function $f(n)$
 - child node is added to frontier if:
 - it has not been reached before
 - re-added if it is now being reached w/ path that has lower path cost than any previous path
 - returns either failure indication or node that represents path to goal

```
function BEST-FIRST-SEARCH(problem, f)  returns solution node or failure
    node ← NODE(STATE=problem.INITIAL)
    frontier ← priority queue ordered by  $f$  , with node as an element
    reached ← lookup table, with one entry with key problem.INITIAL and value node
    while not IS-EMPTY(frontier) do
        node ← POP(frontier)
        if problem.IS-GOAL(node.STATE) then return node
        for each child in EXPAND(problem, node) do
            s ← child.STATE
            if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
                reached[s] ← child
                add child to frontier
    return failure
```

Search Algorithms

Search data structures

- ❑ data structure to store search tree → node represented by data structure with:
 - node.STATE: state to which node corresponds
 - node.PARENT: node in tree that generated this node
 - node.ACTION: action that was applied to parent's state to generate this node
 - node.PATH-COST: total cost of path from initial state to this node, $g(\text{node})$
- ❑ following PARENT pointers back from node → recover states/actions along path
 - Doing this from goal node gives us solution
- ❑ data structure to store frontier → queue of some kind, operations are:
 - IS-EMPTY(frontier) returns true only if there are no nodes in frontier.
 - POP(frontier) removes top node from frontier and returns it.
 - TOP(frontier) returns (but does not remove) top node of frontier.
 - ADD(node, frontier) inserts node into its proper place in queue.
- ❑ candidates: priority queue (best-first), FIFO queue (breadth-first), LIFO queue (depth-first)
- ❑ reached states can be stored as lookup/hash table: key=state , value=node

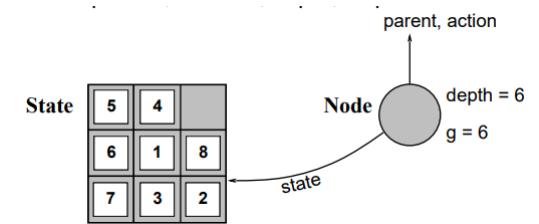
Search Algorithms

Search data structures

Implementation note:

- states vs. nodes:
 - **node:** is data structure constituting part of search tree: includes parent, children, depth, path cost $g(x)$
 - **state:** is a (representation of) a physical configuration: do not have parents, children, depth, or path cost!
- Expand function:
 - creates new nodes
 - filling in various fields
 - use SuccessorFn of problem to create corresponding states

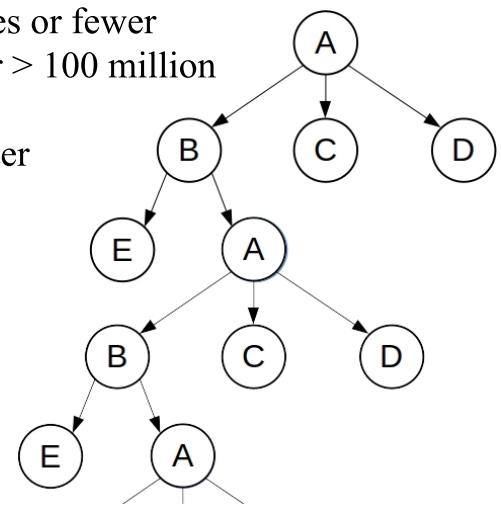
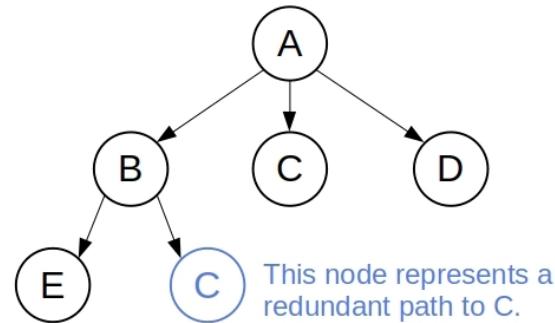
```
function Tree-Search(problem, strategy) returns a solution, or failure
  initialize search tree using initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if node contains a goal state then return corresponding solution
    else expand node and add resulting nodes to search tree
  end
```



Search Algorithms

Redundant paths

- ❑ Repeated state in search tree can be generated by cycles (loopy path)
 - even though state space has finite states, loopy path complete search tree is infinite
 - no limit to how often one can traverse a loop
- ❑ cycle is special case of redundant path
 - Consider agent in 10x10 grid world, with ability to move to any of 8 adjacent squares
 - If there are no obstacles, agent can reach any of the 100 squares in 9 moves or fewer
 - But number of paths of length 9 is almost 8^9 (because of edges of grid), or > 100 million
 - average cell can be reached by over million redundant paths of length 9
 - eliminate redundant paths → complete search roughly a million times faster



Search Algorithms

Redundant paths

- ❑ First, remember all previously reached states → detect all redundant paths
 - **Graph search:** keep only best path to each state, appropriate where many redundant paths
 - preferred choice when table of reached states will fit in memory

```
def depth_first_search(self, value):  
    if self.value == value:  
        return True  
  
    for child in self.children:  
        if child.depth_first_search(value):  
            return True  
    return False
```

```
def depth_first_search(node, value, visited):  
    if node.value == value:  
        return True  
visited.add(node)  
    for neighbor in node.neighbors:  
        if neighbor not in visited:  
            if depth_first_search(neighbor, value, visited):  
                return True  
    return False
```

- ❑ Second, not worry about repeating past: rare/impossible for two paths to reach same state
 - Ex. assembly problem where each action adds part to an evolving assemblage, in some order
 - save memory space if we don't track reached states → don't check for redundant paths
- ❑ Third, compromise and check for cycles, but not for redundant paths in general

Search Algorithms

Measuring problem-solving performance

- ❑ Search strategy: picking order of node expansion
 - evaluated along following dimensions:
 - Completeness: does it always find a solution if one exists?
 - Cost optimality: does it always find a least-cost solution?
 - Time complexity: number of nodes generated/expanded
 - Space complexity: maximum number of nodes in memory
 - Time and space complexity are measured in terms of
 - b: maximum branching factor of the search tree
 - Eight puzzle has effective branching factor of 2.13
 - Rubik's cube has effective branching factor of 13.34
 - Chess has a branching factor of about 35
 - d: depth of the least-cost solution
 - m: maximum depth of the state space (may be ∞)

Search Algorithms

Uninformed search strategies

- ❑ Uninformed strategies use only information available in problem definition
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search

Search Algorithms

Uninformed search: Breadth-first search

Search Algorithms

Uninformed search: Breadth-first search

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - # of actions it takes to reach node

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - # of actions it takes to reach node
- ❑ can get additional efficiency with couple of tricks

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - # of actions it takes to reach node
- ❑ can get additional efficiency with couple of tricks
 - FIFO queue faster than priority queue
→ give correct order of nodes

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - # of actions it takes to reach node
- ❑ can get additional efficiency with couple of tricks
 - FIFO queue faster than priority queue
 - give correct order of nodes
 - new nodes are always deeper than their parents
 - go to back of queue

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - # of actions it takes to reach node
- ❑ can get additional efficiency with couple of tricks
 - FIFO queue faster than priority queue
 - give correct order of nodes
 - new nodes are always deeper than their parents
 - go to back of queue
 - old nodes which are shallower than new nodes
 - get expanded first

https://docs.google.com/presentation/d/e/2PACX-1vRDUQcHg5jUIXPtsrh548S8rEC-1OYj16JOKf3pTNC8jQMod9HnpYV4DiMPO5sMXnDpAJ7bvWjmJle-/embed?start=false&loop=false&delayms=30...

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - # of actions it takes to reach node
- ❑ can get additional efficiency with couple of tricks
 - FIFO queue faster than priority queue
 - give correct order of nodes
 - new nodes are always deeper than their parents
 - go to back of queue
 - old nodes which are shallower than new nodes
 - get expanded first
 - Reached can be states set, rather than state-node map

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - # of actions it takes to reach node
- ❑ can get additional efficiency with couple of tricks
 - FIFO queue faster than priority queue
 - give correct order of nodes
 - new nodes are always deeper than their parents
 - go to back of queue
 - old nodes which are shallower than new nodes
 - get expanded first
 - Reached can be states set, rather than state-node map
 - once a state is reached → never find better path to it
 - do early goal test

Search Algorithms

Uninformed search: Breadth-first search

- ❑ systematic search strategy → complete even on infinite state spaces
 - 1-root is expanded
 - 2-all root successors are expanded next
 - 3-then their successors and so on
- ❑ BEST-FIRST-SEARCH where evaluation function $f(n)$ is node depth
 - o # of actions it takes to reach node
- ❑ can get additional efficiency with couple of tricks
 - o FIFO queue faster than priority queue
 - give correct order of nodes
 - new nodes are always deeper than their parents
 - go to back of queue
 - old nodes which are shallower than new nodes
 - get expanded first
 - o Reached can be states set, rather than state-node map
 - once a state is reached → never find better path to it
 - do early goal test

```

function BREADTH-FIRST-SEARCH(problem) returns solution or failure
  node ← NODE(problem.INITIAL)
  if problem.IS-GOAL(node.STATE) then return node
  frontier ← FIFO queue, with node as element
  reached ← problem.INITIALg
  while not IS-EMPTY(frontier) do
    node ← POP(frontier)
    for each child in EXPAND(problem, node) do
      s ← child.STATE
      if problem.IS-GOAL(s) then return child
      if s is not in reached then
        add s to reached
        add child to frontier
  return failure

```

Search Algorithms

Uninformed search: Breadth-first search

❑ Properties

- Complete? Yes
- Time? $1+b+b^2+b^3+\dots+b^d = O(b^d)$: exp. in d
- Space? $O(b^d)$ (keeps every node in memory)
- Optimal? Yes (if cost = 1/step), not in general

❑ Space is big problem

- typical real-world problem: $b = 10$
processing speed 10^6 nodes/sec
memory requirements of 1 Kbyte/node
- depth $d = 10 \rightarrow$ take less than 3 hours
10 terabytes
- depth $d = 14$, even with infinite memory
takes 3.5 years

Search Algorithms

Uninformed search: Breadth-first search

❑ Properties

- Complete? Yes
- Time? $1+b+b^2+b^3+\dots+b^d = O(b^d)$: exp. in d
- Space? $O(b^d)$ (keeps every node in memory)
- Optimal? Yes (if cost = 1/step), not in general

❑ Space is big problem

- typical real-world problem: $b = 10$
processing speed 10^6 nodes/sec
memory requirements of 1 Kbyte/node
- depth $d = 10 \rightarrow$ take less than 3 hours
10 terabytes
- depth $d = 14$, even with infinite memory
takes 3.5 years

Search Algorithms

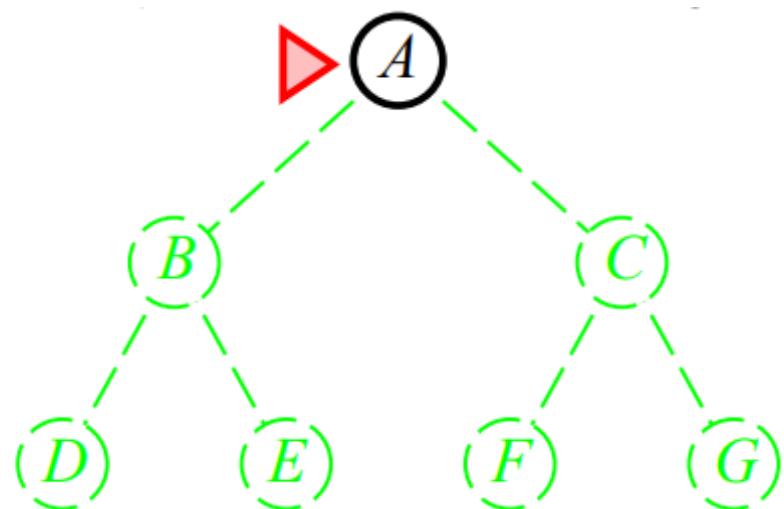
Uninformed search: Breadth-first search

❑ Properties

- Complete? Yes
- Time? $1+b+b^2+b^3+\dots+b^d = O(b^d)$: exp. in d
- Space? $O(b^d)$ (keeps every node in memory)
- Optimal? Yes (if cost = 1/step), not in general

❑ Space is big problem

- typical real-world problem: $b = 10$
processing speed 10^6 nodes/sec
memory requirements of 1 Kbyte/node
- depth $d = 10 \rightarrow$ take less than 3 hours
10 terabytes
- depth $d = 14$, even with infinite memory
takes 3.5 years



Search Algorithms

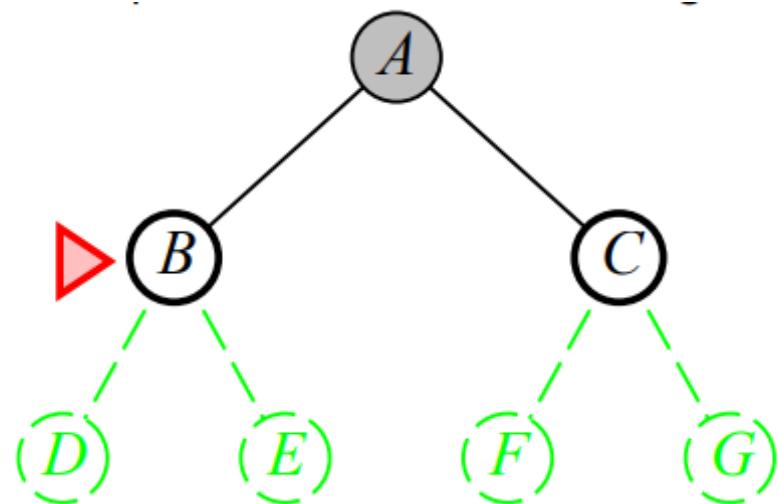
Uninformed search: Breadth-first search

Properties

- Complete? Yes
- Time? $1+b+b^2+b^3+\dots+b^d = O(b^d)$: exp. in d
- Space? $O(b^d)$ (keeps every node in memory)
- Optimal? Yes (if cost = 1/step), not in general

Space is big problem

- typical real-world problem: $b = 10$
processing speed 10^6 nodes/sec
memory requirements of 1 Kbyte/node
- depth $d = 10 \rightarrow$ take less than 3 hours
10 terabytes
- depth $d = 14$, even with infinite memory
takes 3.5 years



Search Algorithms

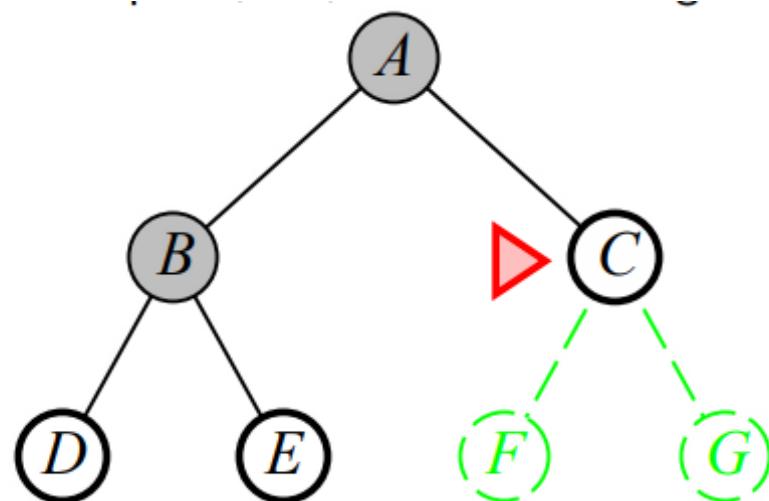
Uninformed search: Breadth-first search

❑ Properties

- Complete? Yes
- Time? $1+b+b^2+b^3+\dots+b^d = O(b^d)$: exp. in d
- Space? $O(b^d)$ (keeps every node in memory)
- Optimal? Yes (if cost = 1/step), not in general

❑ Space is big problem

- typical real-world problem: $b = 10$
processing speed 10^6 nodes/sec
memory requirements of 1 Kbyte/node
- depth $d = 10 \rightarrow$ take less than 3 hours
10 terabytes
- depth $d = 14$, even with infinite memory
takes 3.5 years



Search Algorithms

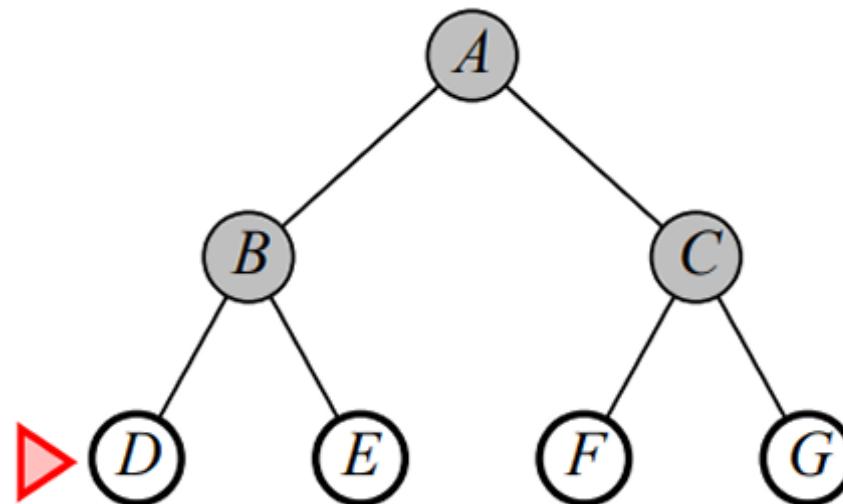
Uninformed search: Breadth-first search

❑ Properties

- Complete? Yes
- Time? $1+b+b^2+b^3+\dots+b^d = O(b^d)$: exp. in d
- Space? $O(b^d)$ (keeps every node in memory)
- Optimal? Yes (if cost = 1/step), not in general

❑ Space is big problem

- typical real-world problem: $b = 10$
processing speed 10^6 nodes/sec
memory requirements of 1 Kbyte/node
- depth $d = 10 \rightarrow$ take less than 3 hours
10 terabytes
- depth $d = 14$, even with infinite memory
takes 3.5 years



Search Algorithms

Uninformed search: Uniform-cost search

Search Algorithms

Uninformed search: Uniform-cost search

Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs

Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node

Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost

Search Algorithms

Uninformed search: Uniform-cost search

- When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node

Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node

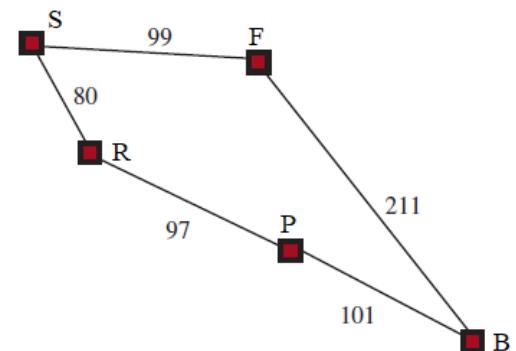
```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

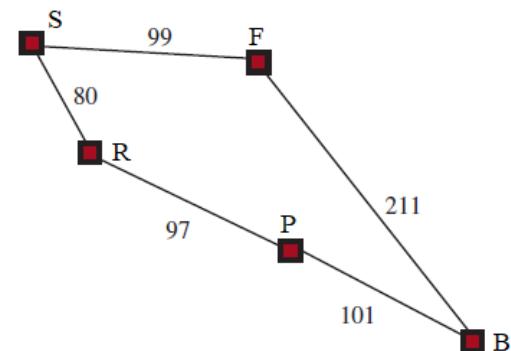


Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node
- ❑ Ex. To reach B from S

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

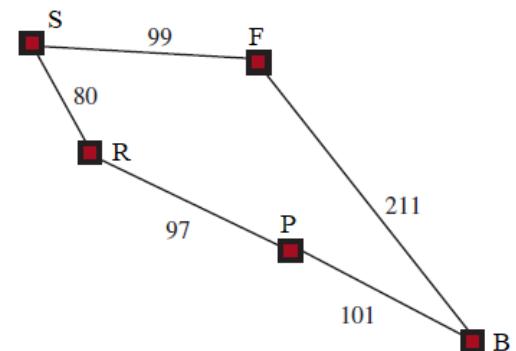


Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node
- ❑ Ex. To reach B from S
 1. POP (S/0), Expand:

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

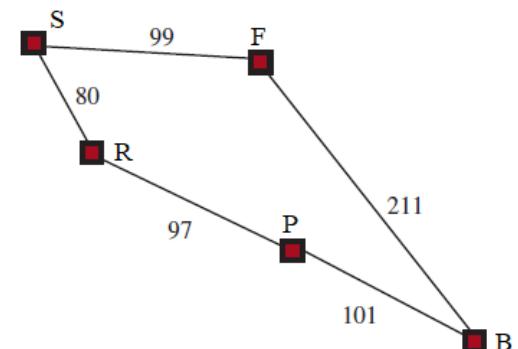


Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node
- ❑ Ex. To reach B from S
 1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

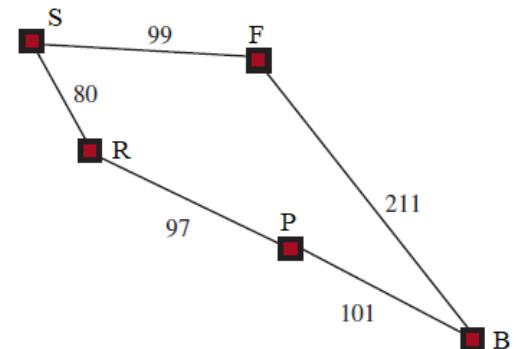


Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node
- ❑ Ex. To reach B from S
 1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

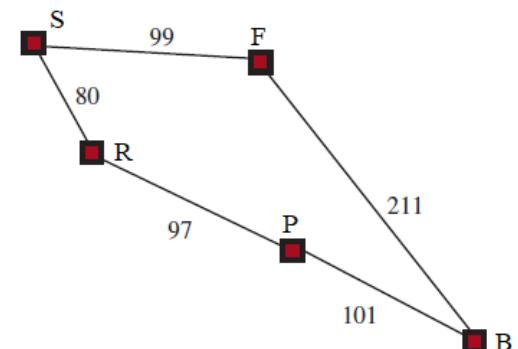


Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node
- ❑ Ex. To reach B from S
 1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
 2. POP (R/80), Expand:

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

When actions have different costs

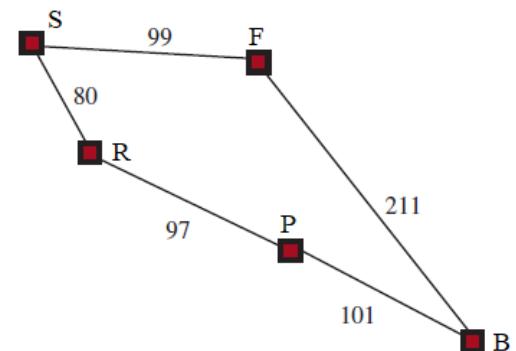
- o obvious choice is use best-first search where evaluation function is cost of path from root to current node
- o idea is search spreads out in waves of uniform path-cost
- o can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

Expand least-cost unexpanded node

Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

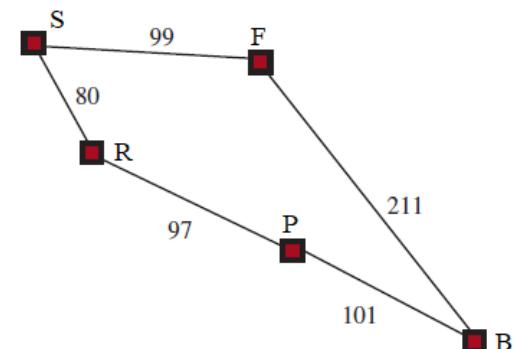


Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node
- ❑ Ex. To reach B from S
 1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
 2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

When actions have different costs

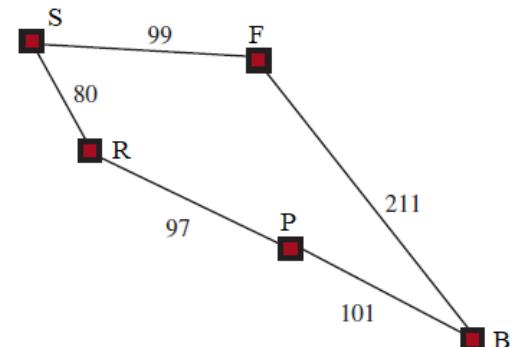
- o obvious choice is use best-first search where evaluation function is cost of path from root to current node
- o idea is search spreads out in waves of uniform path-cost
- o can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

Expand least-cost unexpanded node

Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
3. POP (F/99), Expand:

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

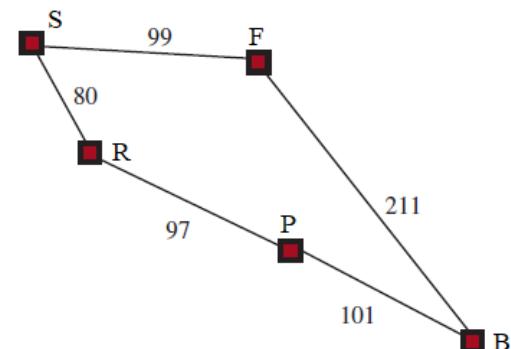
- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

- ❑ Expand least-cost unexpanded node

- ❑ Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
3. POP (F/99), Expand:
 - Successor(s): (B/99+211=310)

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

When actions have different costs

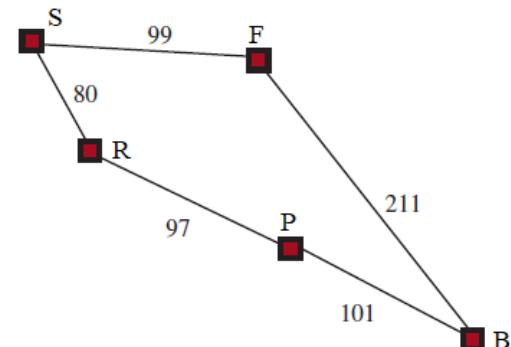
- o obvious choice is use best-first search where evaluation function is cost of path from root to current node
- o idea is search spreads out in waves of uniform path-cost
- o can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

Expand least-cost unexpanded node

Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
3. POP (F/99), Expand:
 - Successor(s): (B/99+211=310)
 - Added to priority queue , Note not detected as goal

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

When actions have different costs

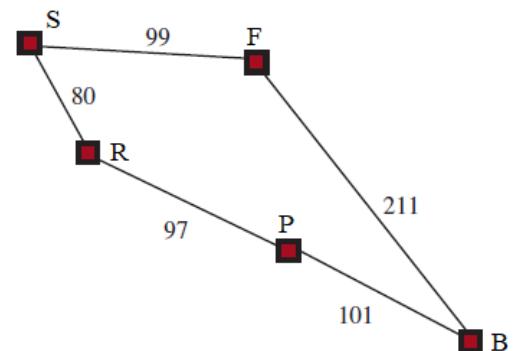
- o obvious choice is use best-first search where evaluation function is cost of path from root to current node
- o idea is search spreads out in waves of uniform path-cost
- o can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

Expand least-cost unexpanded node

Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
3. POP (F/99), Expand:
 - Successor(s): (B/99+211=310)
 - Added to priority queue , Note not detected as goal
4. POP (P/177), Expand:

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

When actions have different costs

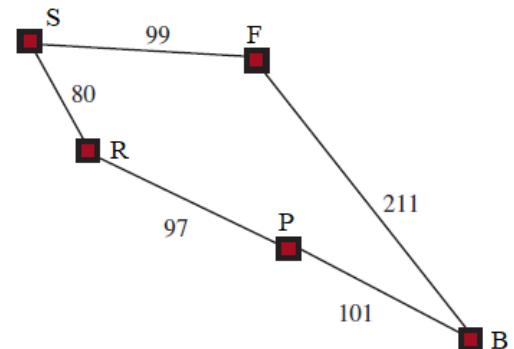
- o obvious choice is use best-first search where evaluation function is cost of path from root to current node
- o idea is search spreads out in waves of uniform path-cost
- o can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

Expand least-cost unexpanded node

Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
3. POP (F/99), Expand:
 - Successor(s): (B/99+211=310)
 - Added to priority queue , Note not detected as goal
4. POP (P/177), Expand:
 - Successor(s): (B/177+101=278)

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```

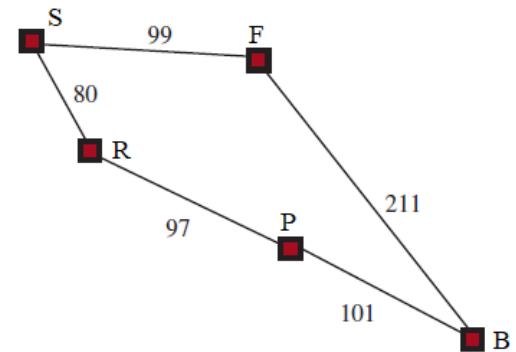


Search Algorithms

Uninformed search: Uniform-cost search

- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function
- ❑ Expand least-cost unexpanded node
- ❑ Ex. To reach B from S
 1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
 2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
 3. POP (F/99), Expand:
 - Successor(s): (B/99+211=310)
 - Added to priority queue , Note not detected as goal
 4. POP (P/177), Expand:
 - Successor(s): (B/177+101=278)
 - Added to priority queue , Note not detected as goal

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

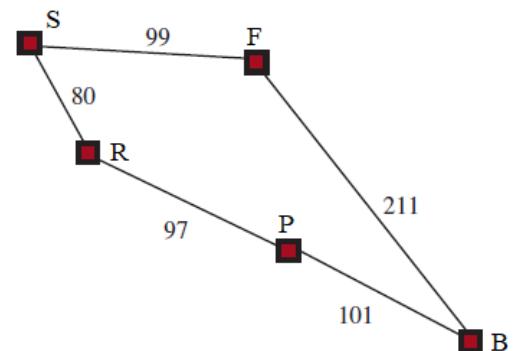
- ❑ When actions have different costs
 - obvious choice is use best-first search where evaluation function is cost of path from root to current node
 - idea is search spreads out in waves of uniform path-cost
 - can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

❑ Expand least-cost unexpanded node

❑ Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
3. POP (F/99), Expand:
 - Successor(s): (B/99+211=310)
 - Added to priority queue , Note not detected as goal
4. POP (P/177), Expand:
 - Successor(s): (B/177+101=278)
 - Added to priority queue , Note not detected as goal
5. POP (B/278)

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

When actions have different costs

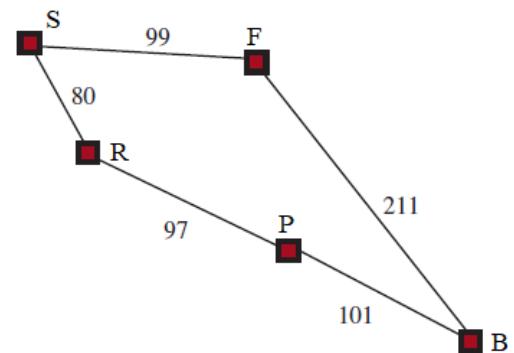
- o obvious choice is use best-first search where evaluation function is cost of path from root to current node
- o idea is search spreads out in waves of uniform path-cost
- o can be implemented as call to BEST-FIRST-SEARCH with PATH-COST as evaluation function

Expand least-cost unexpanded node

Ex. To reach B from S

1. POP (S/0), Expand:
 - Successors of (S): (R/80) , (F/99)
 - Added to priority queue (prioritize lowest cost)
2. POP (R/80), Expand:
 - Successor(s): (P/80+97=177)
 - Added to priority queue
3. POP (F/99), Expand:
 - Successor(s): (B/99+211=310)
 - Added to priority queue , Note not detected as goal
4. POP (P/177), Expand:
 - Successor(s): (B/177+101=278)
 - Added to priority queue , Note not detected as goal
5. POP (B/278)
 - Detect Goal!

```
function UNIFORM-COST-SEARCH(problem)
  returns a solution node, or failure
  return BEST-FIRST-SEARCH(problem,
    PATH-COST)
```



Search Algorithms

Uninformed search: Uniform-cost search

- Implementation:
 - fringe = priority queue ordered by path cost, lowest first
 - Equivalent to breadth-first if step costs all equal
- Properties
 - Complete? Yes, if action cost $\geq \epsilon > 0$
 - where ϵ is lower bound on cost of each action
 - Time? in worst case $O(b^{\lceil C^*/\epsilon \rceil})$
 - where C^* is cost of optimal solution
 - Space? in worst case, $O(b^{\lceil C^*/\epsilon \rceil})$
 - Optimal? Yes, nodes expanded in increasing order of cost

Search Algorithms

Uninformed search: Depth-first search

- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No

Search Algorithms

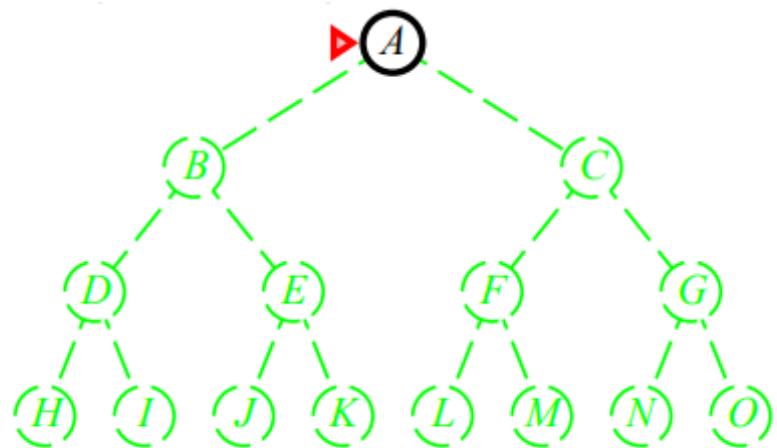
Uninformed search: Depth-first search

- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No

Search Algorithms

Uninformed search: Depth-first search

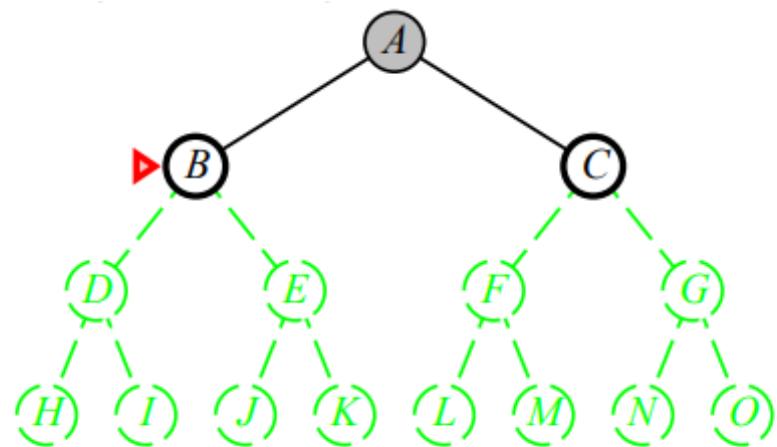
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b.m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

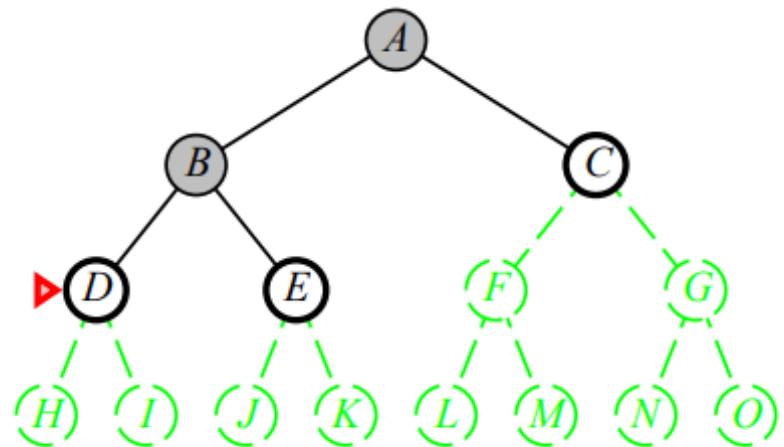
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

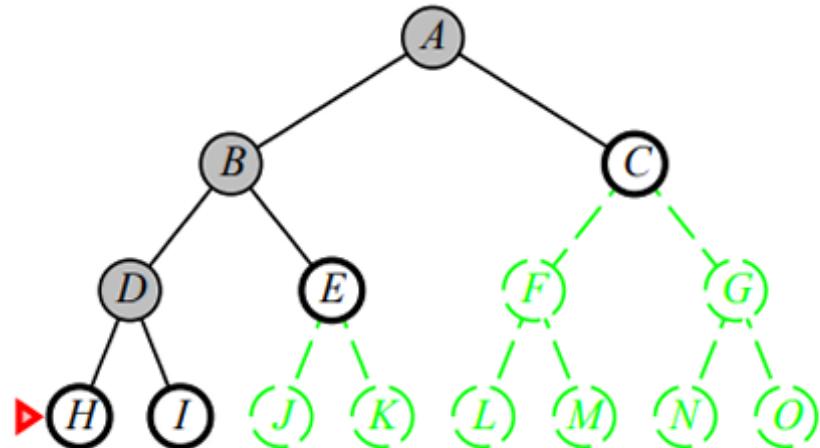
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

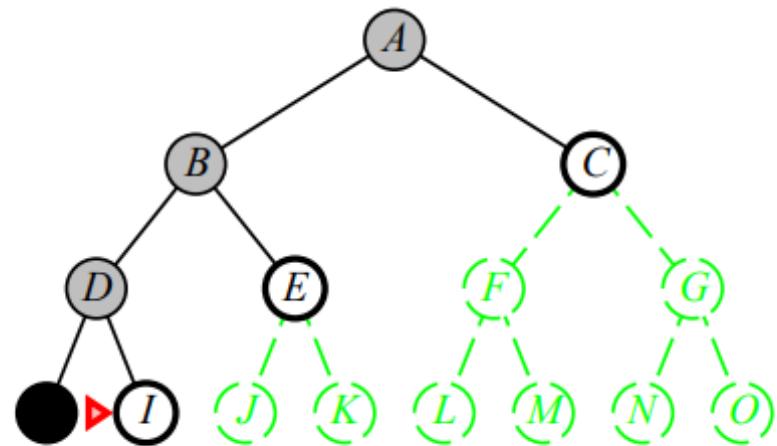
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

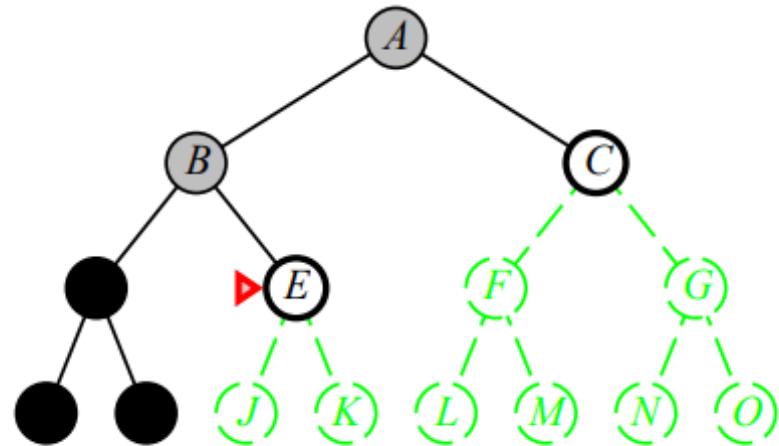
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

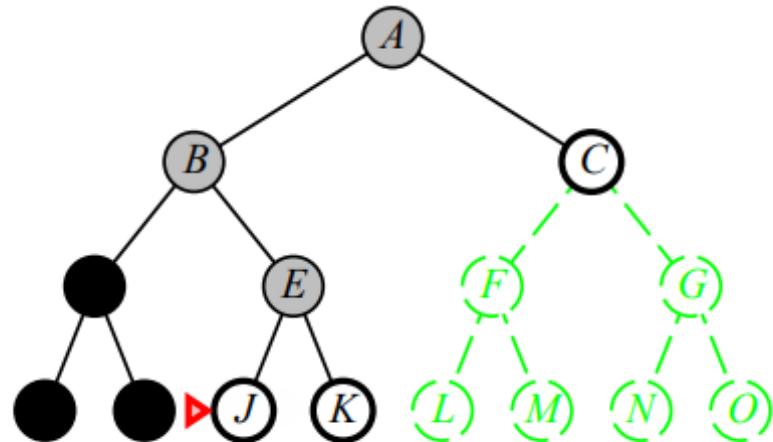
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

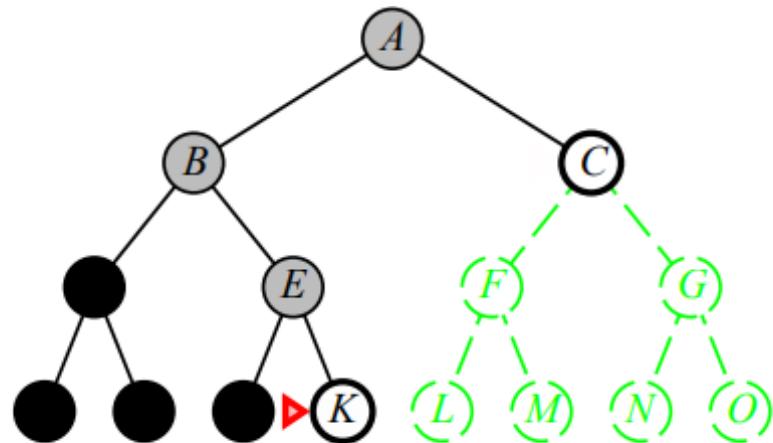
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

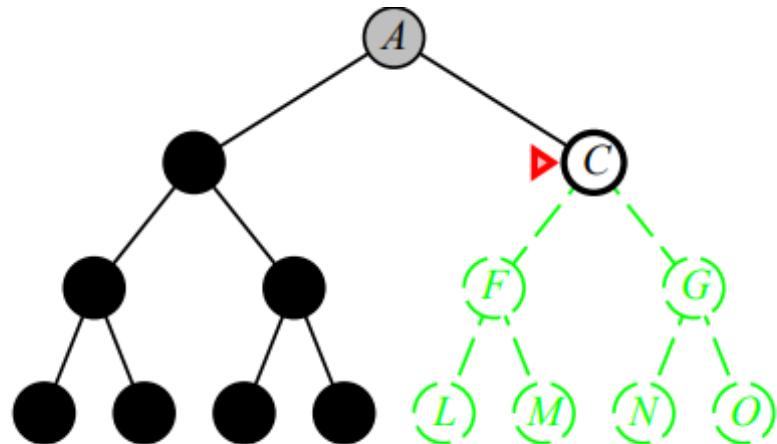
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

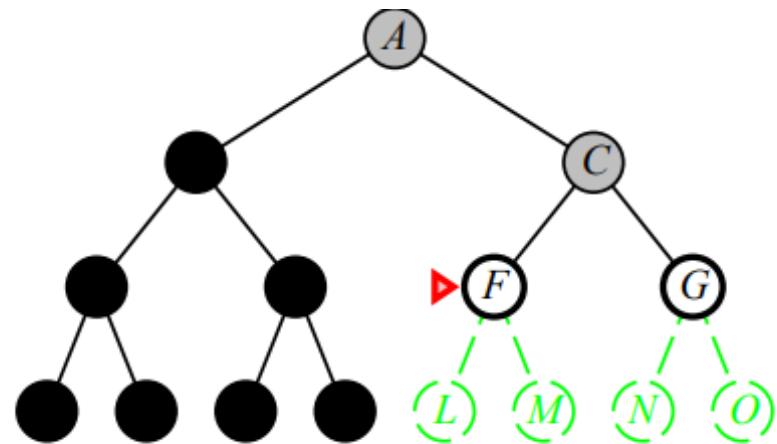
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b.m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

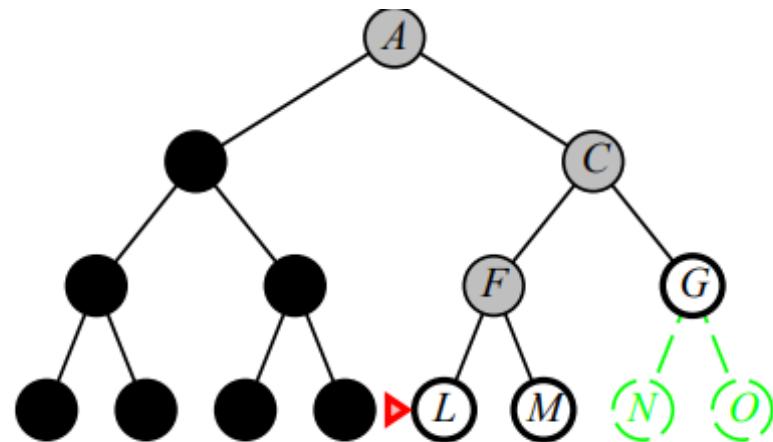
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

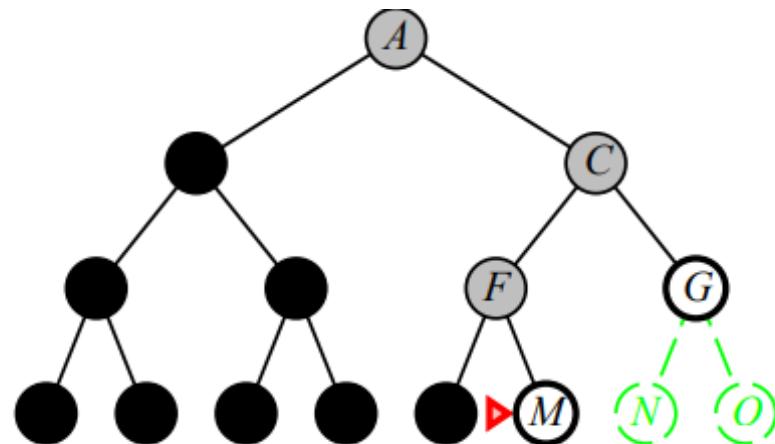
- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-first search

- ❑ Expand deepest unexpanded node
- ❑ Implementation:
 - fringe = LIFO queue, i.e., put successors at front
- ❑ Properties
 - Complete? No: fails in infinite-depth spaces
 - spaces with loops
 - Modify to avoid repeated states along path
→ complete in finite spaces
 - Time? $O(b^m)$: terrible if $m \gg d$
 - but if solutions are dense → may be much faster than breadth-first
 - Space? $O(b \cdot m)$, i.e., linear space!
 - Optimal? No



Search Algorithms

Uninformed search: Depth-limited search

- ❑ Same as depth-first search with depth limit l
 - nodes at depth l have no successors
 - can eliminate cycles at cost of some computation time.
 - If we look only a few links up in parent chain we can catch most cycles
 - good depth limit can be chosen based on knowledge of the problem
 - Diameter of the state-space graph → better depth limit → more efficient depth-limited search

- ❑ Implementation:
 - Recursive implementation

```

function Depth-Limited-Search(problem,limit) returns soln/fail/cutoff
  Recursive-DLS(Make-Node(Initial-State[problem]),problem,limit)

function Recursive-DLS(node,problem,limit) returns soln/fail/cutoff
  cutoff-occurred? ← false
  if Goal-Test(problem, State[node]) then return node
  else if Depth[node] = limit then return cutoff
  else for each successor in Expand(node,problem) do
    result ← Recursive-DLS(successor,problem,limit)
    if result = cutoff then cutoff-occurred? ← true
    else if result ≠ failure then return result
  if cutoff-occurred? then return cutoff else return failure

```

Search Algorithms

Uninformed search: Iterative deepening search

- ❑ Solves problem of picking a good value for l
 - by trying all values: 0, 1, 2, and so on
 - until either solution is found, or depth limited search returns failure value rather than the cutoff value
- ❑ Combines many of the benefits of depth-first and breadth-first search
 - Like DFS: its memory requirements are modest
 - $O(bd)$ when there is solution
 - $O(bm)$ on finite state spaces with no solution
 - Like BFS:
 - optimal for problems where all actions have same cost
 - complete on finite acyclic state spaces, or on any finite state space (check nodes for cycles all way up path)

```
function Iterative-Deepening-Search(problem) returns a solution
    inputs: problem, a problem
    for depth ← 0 to ∞ do
        result ← Depth-Limited-Search( problem, depth)
        if result ≠ cutoff then return result
    end
```

Search Algorithms

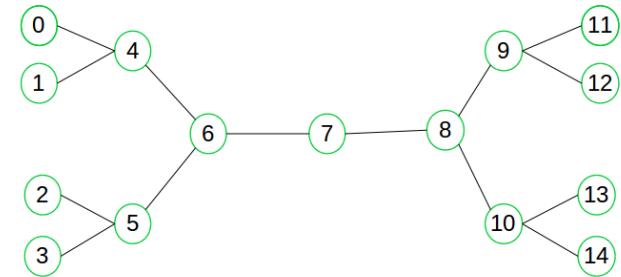
Uninformed search: Bi-directional Search (BDS)

❑ Main idea:

- start searching from both initial state and goal state
 - Forward search from source state toward goal state
 - Backward search from goal state toward source state
- meet in middle.

❑ Properties

- Complete? Yes
- Optimal? Yes
- Time Complexity: $O(b^{d/2})$
- Space Complexity: $O(b^{d/2})$



Search Algorithms

Uninformed search: Comparison

	BFS	UCS	DFS	DLS	IDS	BDS*
Complete?	Yes	Yes	No	If $l \geq d$	Yes	Yes
Optimal?	Yes	Yes	No	No	Yes	Yes
Time	b^d	b^d	b^m	b^l	b^d	$b^{d/2}$
Space	b^d	b^d	$b.m$	$b.l$	$b.d$	$b^{d/2}$

* if applicable, with BDS

d : branching factor

b : depth of solution

m : depth of problem

l : depth limit